

Conditionals: Advanced ORDER BY

Complementary SQL: Solutions

May 11, 2025

Contents

1	Solutions: Practice Meanings, Values, Relations, and Advantages	4
1.1	Solution 1.1: Ordering by Multiple Columns	4
1.2	Solution 1.2: Using NULLS FIRST	4
1.3	Solution 1.3: Using NULLS LAST and Multiple Columns	4
2	Solutions: Practice Disadvantages	5
2.1	Solution 2.1: Disadvantage of Overly Complex Sorting (Readability/Maintainability)	5
2.2	Solution 2.2: Disadvantage of Potentially Misleading Prioritization with NULLS FIRST/LAST	5
3	Solutions: Practice Cases of Lost Advantages	6
3.1	Solution 3.1: Inefficient Simulation of NULLS FIRST	6
3.2	Solution 3.2: Inefficient Custom Sort Order Implementation	6
4	Solution: Hardcore Problem Combining Concepts	8
4.1	Solution 4.1: Comprehensive Employee Ranking and Reporting	8

Dataset for PostgreSQL

The following SQL code creates and populates the necessary tables for the exercises. This dataset is identical to the one provided in the exercises document.

```
1  -- Drop tables if they exist to ensure a clean setup
2  DROP TABLE IF EXISTS employee_projects CASCADE;
3  DROP TABLE IF EXISTS employees CASCADE;
4  DROP TABLE IF EXISTS departments CASCADE;
5
6  -- Create departments table
7  CREATE TABLE departments (
8      department_name VARCHAR(50) PRIMARY KEY,
9      location VARCHAR(50),
10     budget_allocation NUMERIC(12,2)
11 );
12
13 -- Populate departments table
14 INSERT INTO departments (department_name, location, budget_allocation) VALUES
15 ('Engineering', 'New York', 500000.00),
16 ('Marketing', 'San Francisco', 300000.00),
17 ('Sales', 'Chicago', 400000.00),
18 ('Human Resources', 'New York', 250000.00),
19 ('Product', 'San Francisco', 350000.00),
20 ('Support', 'Remote', 150000.00);
21
22 -- Create employees table
23 CREATE TABLE employees (
24     id SERIAL PRIMARY KEY,
25     first_name VARCHAR(50),
26     last_name VARCHAR(50),
27     department VARCHAR(50) REFERENCES departments(department_name),
28     salary NUMERIC(10, 2),
29     hire_date DATE,
30     bonus_percentage NUMERIC(3, 2), -- Nullable
31     manager_id INTEGER REFERENCES employees(id) -- Nullable
32 );
33
34 -- Populate employees table
35 INSERT INTO employees (first_name, last_name, department, salary, hire_date,
36     bonus_percentage, manager_id) VALUES
37 ('Alice', 'Smith', 'Engineering', 90000.00, '2020-03-15', 0.10, NULL),
38 ('Bob', 'Johnson', 'Engineering', 95000.00, '2019-07-01', 0.12, 1),
39 ('Charlie', 'Williams', 'Marketing', 70000.00, '2021-01-10', 0.08, NULL),
40 ('David', 'Brown', 'Sales', 80000.00, '2020-11-05', NULL, NULL), -- Null bonus
41 ('Eve', 'Jones', 'Engineering', 90000.00, '2021-05-20', 0.10, 1),
42 ('Frank', 'Garcia', 'Marketing', 72000.00, '2022-02-01', NULL, 3), -- Null bonus
43 ('Grace', 'Miller', 'Sales', 82000.00, '2019-05-20', 0.09, 4),
44 ('Heidi', 'Davis', 'Human Resources', 65000.00, '2023-01-15', 0.05, NULL),
45 ('Ivan', 'Rodriguez', 'Product', 110000.00, '2020-08-24', 0.15, NULL),
46 ('Judy', 'Martinez', 'Product', 105000.00, '2021-06-10', NULL, 9), -- Null bonus
47 ('Kevin', 'Hernandez', 'Engineering', 88000.00, '2023-03-01', 0.07, 2),
48 ('Linda', 'Lopez', 'Marketing', 68000.00, '2023-04-10', 0.06, 3),
49 ('Mike', 'Gonzalez', 'Sales', 78000.00, '2022-07-18', 0.11, 4),
50 ('Nancy', 'Wilson', 'Human Resources', 67000.00, '2022-09-01', NULL, 8), -- Null bonus
51 ('Olivia', 'Anderson', 'Engineering', 90000.00, '2020-03-15', NULL, 1), -- Null bonus,
52     same salary/hire_date as Alice
53 ('Peter', 'Lee', 'Product', 100000.00, '2021-08-15', 0.12, 9),
54 ('Zoe', 'King', 'Engineering', 92000.00, '2022-05-01', 0.11, 1),
55 ('Yasmin', 'Scott', 'Marketing', 75000.00, '2021-11-20', NULL, 3),
56 ('Eva', 'Taylor', 'Engineering', 90000.00, '2021-05-20', NULL, 1); -- Same salary/
57     hire_date as Eve, but NULL bonus
58
59 -- Create employee_projects table
60 CREATE TABLE employee_projects (
61     employee_id INTEGER REFERENCES employees(id),
62     project_name VARCHAR(100),
63     project_role VARCHAR(50),
64     hours_assigned INTEGER,
65     PRIMARY KEY (employee_id, project_name)
66 );
```

```

65 -- Populate employee_projects table
66 INSERT INTO employee_projects (employee_id, project_name, project_role, hours_assigned)
    VALUES
67 (1, 'Alpha Platform', 'Developer', 120),
68 (1, 'Beta Feature', 'Lead Developer', 80),
69 (2, 'Alpha Platform', 'Senior Developer', 150),
70 (3, 'Campaign X', 'Coordinator', 100),
71 (5, 'Gamma Initiative', 'Developer', 90),
72 (5, 'Alpha Platform', 'Tester', 40),
73 (6, 'Campaign Y', 'Analyst', 110),
74 (7, 'Client Outreach', 'Manager', 60),
75 (9, 'Omega Product', 'Product Owner', 180),
76 (10, 'Omega Product', 'UX Designer', 70),
77 (11, 'Gamma Initiative', 'Junior Developer', 100),
78 (13, 'Client Retention', 'Specialist', 50);

```

Listing 1: Dataset for Advanced ORDER BY Exercises

1 Solutions: Practice Meanings, Values, Relations, and Advantages

1.1 Solution 1.1: Ordering by Multiple Columns

```
1 SELECT first_name, last_name, department, salary
2 FROM employees
3 ORDER BY department ASC, salary DESC;
```

Listing 2: Solution to Exercise 1.1

Problem solved with the given dataset: This query retrieves employee names, their department, and salary from the `employees` table. The results are structured to first group employees by the `department` column, and then, within each department, arrange them from the highest to the lowest `salary`. This demonstrates the advantage of multi-column ordering in creating a more detailed and hierarchical view of the data, allowing for precise control over data presentation.

1.2 Solution 1.2: Using NULLS FIRST

```
1 SELECT first_name, last_name, department, bonus_percentage
2 FROM employees
3 ORDER BY bonus_percentage ASC NULLS FIRST;
```

Listing 3: Solution to Exercise 1.2

Problem solved with the given dataset: This query lists employees along with their department and bonus percentages from the `employees` table. The specific requirement is to prioritize employees with NULL values in the `bonus_percentage` column by showing them first. They are followed by other employees sorted by their `bonus_percentage` from lowest to highest. This highlights the advantage and value of `NULLS FIRST` in explicitly handling and positioning records with missing data at the beginning of the sorted result set, which can be useful for review or special handling.

1.3 Solution 1.3: Using NULLS LAST and Multiple Columns

```
1 SELECT first_name, last_name, hire_date, bonus_percentage
2 FROM employees
3 WHERE department = 'Engineering'
4 ORDER BY hire_date ASC, bonus_percentage DESC NULLS LAST;
```

Listing 4: Solution to Exercise 1.3

Problem solved with the given dataset: This query focuses on employees in the 'Engineering' department. The sorting criteria are: primarily by `hire_date` (earliest first). For ties in `hire_date`, employees who have a value in their `bonus_percentage` column are prioritized (and sorted by the highest bonus first), while those with NULL in `bonus_percentage` are placed at the end of that `hire_date` group. This demonstrates combining multi-column sort with `NULLS LAST` for refined ordering, where records with missing data in a secondary sort key are deprioritized within the primary sort groups.

2 Solutions: Practice Disadvantages

2.1 Solution 2.1: Disadvantage of Overly Complex Sorting (Readability/Maintainability)

```
1 -- This is a conceptual question. The query below is for context.
2 SELECT first_name, last_name, department, hire_date, salary,
   bonus_percentage
3 FROM employees
4 ORDER BY
5     department ASC,
6     (CASE WHEN hire_date < '2021-01-01' THEN 0 ELSE 1 END) ASC, --
   Custom logic for hire_date grouping
7     salary DESC NULLS LAST, -- Salary
   with NULLs at the end
8     (bonus_percentage IS NULL) ASC, --
   Prioritizes non-NULL bonus_percentage
9     last_name ASC;
```

Listing 5: Conceptual Query for Exercise 2.1

Problem illustrated by the example: The provided `ORDER BY` clause sorts employee data from the `employees` table using five distinct criteria, including `CASE` expressions to create synthetic sort keys and explicit `NULLS LAST`. The main disadvantage is significantly reduced readability and maintainability. Such complexity makes the query's intent hard to grasp quickly, increases the likelihood of errors during modification, and makes it difficult for other developers to understand and debug. While powerful, excessive complexity in `ORDER BY` can outweigh the benefits if simpler, more declarative constructs (like direct use of `NULLS FIRST/LAST` where appropriate instead of `CASE` for null handling) could achieve similar results with greater clarity.

2.2 Solution 2.2: Disadvantage of Potentially Misleading Prioritization with `NULLS FIRST/LAST`

```
1 -- This is a conceptual question. The query below is for context.
2 SELECT first_name, last_name, bonus_percentage
3 FROM employees
4 ORDER BY bonus_percentage ASC NULLS FIRST;
```

Listing 6: Conceptual Query for Exercise 2.2

Problem illustrated by the example: The query sorts employees by `bonus_percentage` ascending, placing `NULL` values first. If a large portion of employees have `NULL` for `bonus_percentage` in the `employees` table, these `NULL` records will dominate the top of the report. The disadvantage is that this might obscure the employees with actual low bonus percentages, or it might give undue prominence to records where the bonus is simply unknown or not applicable. If the intent was to focus on those with *defined* low bonuses, `NULLS FIRST` could be counterproductive, making the initial part of the report less informative about the actual data distribution of non-null bonuses. The choice of `NULLS FIRST` vs. `NULLS LAST` must align with the semantic meaning of `NULL` in that context.

3 Solutions: Practice Cases of Lost Advantages

3.1 Solution 3.1: Inefficient Simulation of NULLS FIRST

```
1 -- Less direct / potentially less clear approach using CASE:
2 SELECT first_name, last_name, department, bonus_percentage
3 FROM employees
4 ORDER BY (CASE WHEN bonus_percentage IS NULL THEN 0 ELSE 1 END) ASC,
           bonus_percentage ASC;
5
6 -- Preferred Advanced ORDER BY approach:
7 SELECT first_name, last_name, department, bonus_percentage
8 FROM employees
9 ORDER BY bonus_percentage ASC NULLS FIRST;
```

Listing 7: Solution to Exercise 3.1 - Comparing Approaches

Problem solved by comparing approaches: The first query uses a CASE statement on the `bonus_percentage` column from the `employees` table. It assigns a sort key (0 for NULL, 1 for non-NULL) to bring NULL bonus percentages to the top, then sorts by the actual `bonus_percentage`. The second query achieves the identical result using `bonus_percentage ASC NULLS FIRST`. The NULLS FIRST approach is preferred because it is more concise, more declarative (clearly states the intent of NULL handling), and generally easier to read and understand. It directly leverages a built-in SQL feature designed for this purpose, potentially allowing the database optimizer to handle the sorting more efficiently than a generic CASE expression. This avoids the common mistake of using a more verbose method when a specialized, simpler one exists.

3.2 Solution 3.2: Inefficient Custom Sort Order Implementation

```
1 -- Superior Advanced ORDER BY Approach:
2 SELECT first_name, last_name, department, salary
3 FROM employees
4 ORDER BY
5     CASE department
6         WHEN 'Sales' THEN 1
7         WHEN 'Engineering' THEN 2
8         ELSE 3
9     END ASC,      -- Primary sort: Custom department order
10    department ASC, -- Secondary sort: Alphabetical for 'other'
11    salary DESC;   -- Tertiary sort: Salary within each department
                  group
```

Listing 8: Solution to Exercise 3.2 - Efficient Approach

Problem solved by the efficient approach: The query sorts employees from the `employees` table. It first establishes a custom order for `department` values ('Sales', then 'Engineering', then others) using a CASE expression. If multiple departments fall into the 'ELSE 3' category, they are then sorted alphabetically by `department` name. Finally, within each of these defined department groupings, employees are sorted by `salary` in descending order. This single-query solution is superior to manual multi-query UNION ALL approaches because it is significantly more efficient (requires only one scan or pass over the data).

by the database engine), is much cleaner to write and maintain, and directly utilizes the power of the `ORDER BY` clause for complex sorting logic. This prevents the loss of performance and clarity that comes from avoiding advanced `ORDER BY` capabilities.

4 Solution: Hardcore Problem Combining Concepts

4.1 Solution 4.1: Comprehensive Employee Ranking and Reporting

```
1 WITH DepartmentTotalHours AS (  
2     -- Calculate total project hours per department  
3     -- Uses: Aggregation (SUM), GROUP BY, HAVING, LEFT JOIN (to include  
4     employees/departments even if no projects assigned to specific  
5     employees)  
6     SELECT  
7         e.department,  
8         SUM(COALESCE(ep.hours_assigned, 0)) AS total_department_hours  
9     FROM employees e  
10    LEFT JOIN employee_projects ep ON e.id = ep.employee_id  
11    GROUP BY e.department  
12    HAVING SUM(COALESCE(ep.hours_assigned, 0)) > 200  
13 ),  
14 RankedEmployees AS (  
15     -- Select and rank employees based on criteria  
16     -- Uses: Joins (INNER JOIN with departments), WHERE (date, list,  
17     subquery),  
18     -- Window Function (RANK() with PARTITION BY and complex ORDER BY  
19     including NULLS LAST),  
20     -- String concatenation  
21     SELECT  
22         e.first_name || ' ' || e.last_name AS full_name,  
23         e.department,  
24         d.location,  
25         e.salary,  
26         e.hire_date,  
27         e.bonus_percentage,  
28         RANK() OVER (  
29             PARTITION BY e.department  
30             ORDER BY  
31                 e.salary DESC,  
32                 e.bonus_percentage DESC NULLS LAST, -- Key: Handles tie  
33                 e.hire_date ASC  
34             ) AS department_rank  
35     FROM employees e  
36     JOIN departments d ON e.department = d.department_name  
37     WHERE e.hire_date >= '2021-01-01' --  
38     Basic SQL: WHERE condition  
39     AND d.location IN ('New York', 'San Francisco') --  
40     Basic SQL: WHERE IN  
41     AND e.department IN (SELECT department FROM DepartmentTotalHours)  
42     -- Complementary SQL: Subquery in WHERE  
43 )  
44 -- Final selection and ordering  
45 -- Uses: COALESCE (Null Space), final ORDER BY (multi-column)  
46 SELECT  
47     full_name,  
48     department,  
49     salary,  
50     hire_date,  
51     COALESCE(bonus_percentage::TEXT, 'N/A') AS bonus_percentage_display
```



```

45     , -- Intermediate SQL: COALESCE, Casters
      department_rank
46 FROM RankedEmployees
47 ORDER BY department ASC, department_rank ASC; -- Basic and Advanced
      ORDER BY

```

Listing 9: Solution to Exercise 4.1

Problem solved with the given dataset: This query addresses a complex reporting requirement by integrating multiple SQL concepts.

1. It first identifies **departments** that are active enough (total project hours \geq 200) using a Common Table Expression (CTE **DepartmentTotalHours**) that joins **employees** and **employee_projects**, aggregates **hours_assigned** using **SUM()**, and filters with **HAVING**. **COALESCE** handles employees with no projects.
2. The main CTE (**RankedEmployees**) then filters **employees** based on **hire_date** (from **employees**), **location** (from **departments** via a **JOIN**), and membership in the previously identified active departments (using a subquery with **IN**).
3. Crucially, for these filtered employees, it calculates a **department_rank** using the **RANK()** window function. This rank is partitioned by **department**. The **ORDER BY** clause within **RANK()** implements the specified multi-level tie-breaking logic: **salary DESC** first, then **bonus_percentage DESC NULLS LAST** (which correctly places non-NULL bonuses first, then sorts them descending, and places all NULL bonuses after non-NULL ones), and finally **hire_date ASC**. This demonstrates a sophisticated use of "Advanced ORDER BY" concepts (**NULLS LAST**, multi-column ordering) within an analytical function.
4. The final **SELECT** statement formats the **bonus_percentage** for display (using **COALESCE** and **CAST**) and orders the entire result set by **department** and then by the calculated **department_rank**.

This exercise showcases how "Advanced ORDER BY" features are essential not only for final result sorting but also within window functions to achieve complex analytical rankings, combined with various other SQL techniques from basic to intermediate levels. The dataset, with its nullable **bonus_percentage**, varied salaries, hire dates, and linked project data, provides the necessary complexity to make such a problem meaningful.