

Date Functions, Cases, and Null Space

Complementary SQL: Solutions

May 13, 2025

Contents

1	Date Functions (Complementary SQL) - Solutions	5
2	Cases (Complementary SQL) - Solutions	9
3	Null Space (Complementary SQL) - Solutions	13
4	Hardcore Problem - Solution	16

Global Dataset for PostgreSQL

The following SQL code creates and populates the necessary tables for the exercises. This dataset should be loaded into your PostgreSQL environment before running the solutions.

```
1  -- Drop tables if they exist to ensure a clean setup
2  DROP TABLE IF EXISTS leave_requests CASCADE;
3  DROP TABLE IF EXISTS employee_projects CASCADE;
4  DROP TABLE IF EXISTS projects CASCADE;
5  DROP TABLE IF EXISTS employees CASCADE;
6  DROP TABLE IF EXISTS departments CASCADE;
7
8  -- Departments Table
9  CREATE TABLE departments (
10     dept_id SERIAL PRIMARY KEY,
11     dept_name VARCHAR(100) NOT NULL UNIQUE,
12     creation_date DATE NOT NULL,
13     location VARCHAR(50)
14 );
15
16 -- Employees Table
17 CREATE TABLE employees (
18     emp_id SERIAL PRIMARY KEY,
19     emp_name VARCHAR(100) NOT NULL,
20     hire_date DATE NOT NULL,
21     salary NUMERIC(10, 2),
22     dept_id INT REFERENCES departments(dept_id),
23     manager_id INT REFERENCES employees(emp_id), -- Self-reference for manager
24     termination_date DATE, -- NULL if currently employed
25     email VARCHAR(100) UNIQUE,
26     performance_rating INT CHECK (performance_rating BETWEEN 1 AND 5 OR
27     performance_rating IS NULL)
28 );
29
30 -- Projects Table
31 CREATE TABLE projects (
32     project_id SERIAL PRIMARY KEY,
33     project_name VARCHAR(100) NOT NULL UNIQUE,
34     start_date DATE NOT NULL,
35     planned_end_date DATE NOT NULL,
36     actual_end_date DATE, -- NULL if not completed
37     budget NUMERIC(12, 2),
38     lead_emp_id INT REFERENCES employees(emp_id),
39     CONSTRAINT check_project_dates CHECK (planned_end_date >= start_date)
40 );
41
42 -- Employee_Projects Table (Junction table)
43 CREATE TABLE employee_projects (
44     emp_project_id SERIAL PRIMARY KEY,
45     emp_id INT REFERENCES employees(emp_id) ON DELETE CASCADE,
46     project_id INT REFERENCES projects(project_id) ON DELETE CASCADE,
47     assigned_date DATE NOT NULL,
48     role VARCHAR(50),
49     hours_billed NUMERIC(6, 2) DEFAULT 0.00,
50     billing_rate NUMERIC(8, 2),
51     completion_date DATE, -- Date employee completed their part
52     UNIQUE(emp_id, project_id) -- An employee has one role per project
53 );
54
55 -- Leave_Requests Table
56 CREATE TABLE leave_requests (
57     leave_id SERIAL PRIMARY KEY,
58     emp_id INT REFERENCES employees(emp_id) ON DELETE CASCADE,
59     request_date TIMESTAMP WITHOUT TIME ZONE DEFAULT CURRENT_TIMESTAMP,
60     leave_start_date DATE NOT NULL,
61     leave_end_date DATE NOT NULL,
62     status VARCHAR(20) CHECK (status IN ('Pending', 'Approved', 'Rejected', 'Cancelled')
63     ),
64     approved_by_manager_id INT REFERENCES employees(emp_id),
65     reason TEXT,
66     CONSTRAINT check_leave_dates CHECK (leave_end_date >= leave_start_date)
67 );
```

```

66
67 -- Populate Departments
68 INSERT INTO departments (dept_name, creation_date, location) VALUES
69 ('Human Resources', '2010-01-15', 'New York'),
70 ('Engineering', '2010-03-01', 'San Francisco'),
71 ('Sales', '2010-02-01', 'Chicago'),
72 ('Marketing', '2011-05-20', 'New York'),
73 ('Finance', '2010-01-20', 'New York');
74
75 -- Populate Employees (managers first, then their reports)
76 INSERT INTO employees (emp_name, hire_date, salary, dept_id, manager_id,
77   termination_date, email, performance_rating) VALUES
78 ('Alice Wonderland', '2015-06-01', 90000.00, 2, NULL, NULL, 'alice@example.com', 5), --
79   Eng Manager
80 ('David Copperfield', '2015-03-10', 120000.00, 3, NULL, NULL, 'david@example.com', 4),
81   -- Sales Manager
82 ('Frankenstein Monster', '2019-11-01', 95000.00, 1, NULL, NULL, 'frank@example.com', 3),
83   -- HR Manager
84 ('Ivy Poison', '2022-09-15', 110000.00, 5, NULL, NULL, 'ivy@example.com', 4); -- Finance
85   Head
86
87 INSERT INTO employees (emp_name, hire_date, salary, dept_id, manager_id,
88   termination_date, email, performance_rating) VALUES
89 ('Bob The Builder', '2016-08-15', 75000.00, 2, (SELECT emp_id from employees WHERE email
90   = 'alice@example.com'), NULL, 'bob@example.com', 4),
91 ('Carol Danvers', '2017-01-20', 80000.00, 2, (SELECT emp_id from employees WHERE email=
92   'alice@example.com'), NULL, 'carol@example.com', 5),
93 ('Eve Harrington', '2018-07-01', 65000.00, 3, (SELECT emp_id from employees WHERE email=
94   'david@example.com'), '2023-12-31', 'eve@example.com', 2),
95 ('Grace Hopper', '2020-02-10', 70000.00, 1, (SELECT emp_id from employees WHERE email='
96   frank@example.com'), NULL, 'grace@example.com', NULL),
97 ('Henry Jekyll', '2021-05-01', 50000.00, 4, (SELECT emp_id from employees WHERE email='
98   david@example.com'), NULL, 'henry@example.com', 3), -- Marketing, reports to Sales
99   head for now
100 ('Jack Sparrow', '2023-01-20', 60000.00, 5, (SELECT emp_id from employees WHERE email='
101   ivy@example.com'), '2023-08-15', 'jack@example.com', 1),
102 ('Kevin McCallister', '2018-09-01', 72000.00, 2, (SELECT emp_id from employees WHERE
103   email='alice@example.com'), NULL, 'kevin@example.com', NULL),
104 ('Laura Croft', '2019-04-15', 85000.00, 4, (SELECT emp_id from employees WHERE email='
105   david@example.com'), NULL, 'laura@example.com', 5); -- Marketing, reports to Sales
106   head
107
108 -- Populate Projects
109 INSERT INTO projects (project_name, start_date, planned_end_date, actual_end_date,
110   budget, lead_emp_id) VALUES
111 ('Project Alpha', '2023-01-01', '2023-06-30', '2023-07-15', 100000.00, (SELECT emp_id
112   from employees WHERE email='alice@example.com')),
113 ('Project Beta', '2023-03-01', '2023-09-30', NULL, 150000.00, (SELECT emp_id from
114   employees WHERE email='bob@example.com')),
115 ('Project Gamma', '2022-09-01', '2023-03-31', '2023-03-20', 80000.00, (SELECT emp_id
116   from employees WHERE email='alice@example.com')),
117 ('Project Delta', '2023-08-01', '2024-02-29', NULL, 200000.00, (SELECT emp_id from
118   employees WHERE email='david@example.com')),
119 ('Project Epsilon', '2023-05-01', '2023-08-31', '2023-09-05', 60000.00, NULL),
120 ('Project Zeta', '2024-01-01', '2024-01-30', NULL, 120000.00, (SELECT emp_id from
121   employees WHERE email='bob@example.com')), -- Ends Jan 30, 2024
122 ('Project Omega', '2023-10-01', '2023-12-20', '2023-12-20', 50000.00, (SELECT emp_id
123   from employees WHERE email='carol@example.com')),
124 ('Critical Eng Proj 1', '2024-01-02', '2024-01-25', NULL, 5000, (SELECT emp_id from
125   employees WHERE email='alice@example.com')), -- Critical for report
126 ('Critical Eng Proj 2', '2024-01-05', '2024-02-10', NULL, 5000, (SELECT emp_id from
127   employees WHERE email='bob@example.com')); -- Critical for report
128
129 -- Populate Employee Projects
130 INSERT INTO employee_projects (emp_id, project_id, assigned_date, role, hours_billed,
131   billing_rate, completion_date) VALUES
132 ((SELECT emp_id from employees WHERE email='alice@example.com'), (SELECT project_id from
133   projects WHERE project_name='Project Alpha'), '2023-01-01', 'Project Manager', 200,
134   150.00, '2023-07-15'),
135 ((SELECT emp_id from employees WHERE email='alice@example.com'), (SELECT project_id from
136   projects WHERE project_name='Project Gamma'), '2022-09-01', 'Project Manager', 180,
137   150.00, '2023-03-20'),

```

```

108 ((SELECT emp_id from employees WHERE email='bob@example.com'), (SELECT project_id from
    projects WHERE project_name='Project Alpha'), '2023-01-05', 'Developer', 300,
    120.00, '2023-07-10'),
109 ((SELECT emp_id from employees WHERE email='bob@example.com'), (SELECT project_id from
    projects WHERE project_name='Project Beta'), '2023-03-01', 'Lead Developer', 250,
    130.00, NULL),
110 ((SELECT emp_id from employees WHERE email='bob@example.com'), (SELECT project_id from
    projects WHERE project_name='Project Zeta'), '2024-01-01', 'Lead Developer', 50,
    135.00, NULL),
111 ((SELECT emp_id from employees WHERE email='carol@example.com'), (SELECT project_id from
    projects WHERE project_name='Project Beta'), '2023-03-05', 'Developer', 220,
    120.00, NULL),
112 ((SELECT emp_id from employees WHERE email='carol@example.com'), (SELECT project_id from
    projects WHERE project_name='Project Epsilon'), '2023-05-01', 'Consultant', 100,
    200.00, '2023-09-05'),
113 ((SELECT emp_id from employees WHERE email='carol@example.com'), (SELECT project_id from
    projects WHERE project_name='Project Omega'), '2023-10-01', 'Developer', 80,
    125.00, '2023-12-20'),
114 ((SELECT emp_id from employees WHERE email='eve@example.com'), (SELECT project_id from
    projects WHERE project_name='Project Delta'), '2023-08-01', 'Sales Rep', 150,
    100.00, '2023-12-31'),
115 ((SELECT emp_id from employees WHERE email='grace@example.com'), (SELECT project_id from
    projects WHERE project_name='Project Alpha'), '2023-02-01', 'HR Coordinator', 80,
    NULL, '2023-07-15'),
116 ((SELECT emp_id from employees WHERE email='kevin@example.com'), (SELECT project_id from
    projects WHERE project_name='Project Beta'), '2023-03-15', 'QA Engineer', 180,
    110.00, NULL),
117 ((SELECT emp_id from employees WHERE email='laura@example.com'), (SELECT project_id from
    projects WHERE project_name='Project Delta'), '2023-08-05', 'Marketing Lead', 200,
    140.00, NULL);
118
119 -- Populate Leave_Requests
120 INSERT INTO leave_requests (emp_id, request_date, leave_start_date, leave_end_date,
    status, approved_by_manager_id, reason) VALUES
121 ((SELECT emp_id from employees WHERE email='bob@example.com'), '2023-04-01 10:00:00', '
    2023-04-10', '2023-04-12', 'Approved', (SELECT emp_id from employees WHERE email='
    alice@example.com'), 'Vacation'),
122 ((SELECT emp_id from employees WHERE email='carol@example.com'), '2023-05-10 14:30:00',
    '2023-06-01', '2023-06-05', 'Approved', (SELECT emp_id from employees WHERE email='
    alice@example.com'), 'Personal Leave'),
123 ((SELECT emp_id from employees WHERE email='eve@example.com'), '2023-11-01 09:00:00', '
    2023-11-10', '2023-11-15', 'Pending', (SELECT emp_id from employees WHERE email='
    david@example.com'), 'Sick Leave'),
124 ((SELECT emp_id from employees WHERE email='grace@example.com'), '2023-06-15 11:00:00',
    '2023-07-01', '2023-07-03', 'Rejected', (SELECT emp_id from employees WHERE email='
    frank@example.com'), NULL),
125 ((SELECT emp_id from employees WHERE email='bob@example.com'), '2023-08-01 16:00:00', '
    2023-08-20', '2023-08-25', 'Approved', (SELECT emp_id from employees WHERE email='
    alice@example.com'), 'Family event'),
126 ((SELECT emp_id from employees WHERE email='kevin@example.com'), '2023-09-01 08:00:00',
    '2023-09-10', '2023-09-11', 'Pending', (SELECT emp_id from employees WHERE email='
    alice@example.com'), ''),
127 ((SELECT emp_id from employees WHERE email='laura@example.com'), '2024-01-10 10:00:00',
    '2024-02-01', '2024-02-05', 'Pending', (SELECT emp_id from employees WHERE email='
    david@example.com'), 'Conference'),
128 ((SELECT emp_id from employees WHERE email='alice@example.com'), '2023-12-01 09:00:00',
    '2023-12-20', '2023-12-28', 'Approved', NULL, 'Holiday Season'),
129 ((SELECT emp_id from employees WHERE email='bob@example.com'), '2024-01-10 10:00:00', '
    2024-01-14', '2024-01-16', 'Approved', (SELECT emp_id from employees WHERE email='
    alice@example.com'), 'Short break'); -- Bob on leave on Jan 15 2024

```

Listing 1: Global Dataset for Exercises

1 Date Functions (Complementary SQL) - Solutions

Concepts: Date arithmetic, OVERLAPS operator.

(i) Meaning, values, relations, advantages of unique usage

Exercise 1.1: Project Timeline Extension and Next Month Check

```
1 SELECT
2     project_name ,
3     planned_end_date AS original_planned_end ,
4     planned_end_date + INTERVAL '2 months' + INTERVAL '15 days' AS
    extended_planned_end ,
5     (EXTRACT(YEAR FROM planned_end_date) = 2024 AND EXTRACT(MONTH FROM
    planned_end_date) = 2) AS is_planned_feb_2024
6 FROM
7     projects ;
```

Listing 2: Solution for Exercise 1.1

Exercise 1.2: Identifying Concurrent Project Assignments for Employees

```
1 SELECT
2     e.emp_name ,
3     p1.project_name AS project1_name ,
4     ep1.assigned_date AS p1_assigned ,
5     ep1.completion_date AS p1_completed ,
6     p2.project_name AS project2_name ,
7     ep2.assigned_date AS p2_assigned ,
8     ep2.completion_date AS p2_completed
9 FROM
10     employee_projects ep1
11 JOIN
12     employees e ON ep1.emp_id = e.emp_id
13 JOIN
14     projects p1 ON ep1.project_id = p1.project_id
15 JOIN
16     employee_projects ep2 ON ep1.emp_id = ep2.emp_id AND ep1.project_id
    < ep2.project_id -- Different projects for the same employee
17 JOIN
18     projects p2 ON ep2.project_id = p2.project_id
19 WHERE
20     (ep1.assigned_date , COALESCE(ep1.completion_date , DATE '9999-12-31'
    )) OVERLAPS
21     (ep2.assigned_date , COALESCE(ep2.completion_date , DATE '9999-12-31'
    ));
```

Listing 3: Solution for Exercise 1.2

(ii) Disadvantages of all its technical concepts

Exercise 1.3: OVERLAPS with identical start/end points *Explanation:* In PostgreSQL, when using (date, date) with OVERLAPS where both dates are identical, this defines an empty interval because the end is not after the start. For example, (DATE '2023-01-01', DATE '2023-01-01') OVERLAPS (DATE '2023-01-01', DATE '2023-01-05') will be FALSE. This is because an empty interval cannot overlap anything. If the intention

is to check if a specific **day** falls within a period, you should ensure the "day" is represented as a one-day interval, e.g., (`start_date, start_date + INTERVAL '1 day'`) if your periods are exclusive of the end date, or ensure the comparison handles points correctly. For checking if a single date falls within a range [S, E], it's usually clearer and more direct to use: `single_date >= S AND single_date <= E`.

```

1  -- Example: Project starts on a day an employee is on approved leave.
2  -- Incorrectly using (start_date, start_date) for OVERLAPS
3  SELECT
4      p.project_name,
5      p.start_date,
6      e.emp_name,
7      lr.leave_start_date,
8      lr.leave_end_date,
9      (p.start_date, p.start_date) OVERLAPS (lr.leave_start_date, lr.
      leave_end_date) AS overlaps_test_empty_interval -- Likely FALSE
10 FROM projects p
11 JOIN employee_projects ep ON p.project_id = ep.project_id
12 JOIN employees e ON ep.emp_id = e.emp_id
13 JOIN leave_requests lr ON e.emp_id = lr.emp_id
14 WHERE lr.status = 'Approved' AND p.project_name = 'Project Alpha' AND e
      .emp_name = 'Alice Wonderland' LIMIT 1; -- Example with data
15
16 -- Correct way to check if a single date is within a range (inclusive
      end date for leave)
17 SELECT
18     p.project_name,
19     p.start_date,
20     e.emp_name,
21     lr.leave_start_date,
22     lr.leave_end_date,
23     p.start_date >= lr.leave_start_date AND p.start_date <= lr.
      leave_end_date AS day_is_within_leave
24 FROM projects p
25 JOIN employee_projects ep ON p.project_id = ep.project_id
26 JOIN employees e ON ep.emp_id = e.emp_id
27 JOIN leave_requests lr ON e.emp_id = lr.emp_id
28 WHERE lr.status = 'Approved' AND p.project_name = 'Project Alpha' AND e
      .emp_name = 'Alice Wonderland' LIMIT 1;

```

Listing 4: Solution for Exercise 1.3 - Demonstrating OVERLAPS and correct check

Disadvantage Summary: The main disadvantage is potential misinterpretation of OVERLAPS with point-in-time events or zero-duration intervals if the database's specific handling (e.g., empty intervals) isn't understood. Direct comparisons are often clearer for single point-in-range checks.

Exercise 1.4: Time Zone Issues in Date Arithmetic without Explicit Time Zone Handling *Discussion of Disadvantage:*

1. **TIMESTAMP WITHOUT TIME ZONE:** If `CURRENT_TIMESTAMP` returns this type, it represents a local time. Adding an interval is straightforward mathematically but the resulting timestamp's meaning is relative to an *assumed* time zone, often the server's. If data or users are in different time zones, this can lead to off-by-several-hours errors when interpreting these timestamps.
2. **TIMESTAMP WITH TIME ZONE:** This type stores UTC or a timestamp with zone information. `CURRENT_TIMESTAMP` usually returns this. Arithmetic is generally con-

sistent (e.g., adding 24 hours). However, when displaying or converting this to local times, issues arise if the target local time zone isn't specified or correctly applied. An interval like `INTERVAL '1 day'` means exactly 24 hours. Adding `INTERVAL '1 day'` across a Daylight Saving Time (DST) transition might result in a local time that isn't the "same time next day" but is 23 or 25 hours later in local wall-clock time.

3. **Ambiguity of "Day"**: `INTERVAL '1 day'` typically means 24 hours. If you need "next calendar day at the same local time", especially across DST, more complex logic involving `DATE_TRUNC`, local time zone conversions, and then adding intervals might be needed.

The disadvantage is that not being mindful of time zones when performing date/timestamp arithmetic can lead to data that is misinterpreted or incorrectly processed.

```

1 -- Example (behavior depends on session time zone for display)
2 SELECT
3     TIMESTAMP WITH TIME ZONE '2024-03-10 01:00:00 America/New_York' AS
    ts_before_dst,
4     TIMESTAMP WITH TIME ZONE '2024-03-10 01:00:00 America/New_York' +
    INTERVAL '24 hours' AS ts_plus_24h;
5     -- In many parts of US, March 10 2024 is a DST spring forward. 2 AM
6     -- becomes 3 AM.
7     -- 1:00 AM EST + 24 hours = 2:00 AM EDT the next day (which is 24
8     -- actual hours later).
9
10 -- For TIMESTAMP WITHOUT TIME ZONE (assuming server time is local and
11 -- not UTC)
12 SELECT
13     '2024-01-15 10:00:00'::TIMESTAMP WITHOUT TIME ZONE AS
    current_ts_no_tz,
14     ('2024-01-15 10:00:00'::TIMESTAMP WITHOUT TIME ZONE) + INTERVAL '1
15     day' AS next_day_no_tz;
16     -- Result: 2024-01-16 10:00:00. Meaning is relative to unspecified
17     -- time zone.

```

Listing 5: Solution for Exercise 1.4 - Conceptual Examples

(iii) Practice entirely cases where people in general does not use these approaches losing their advantages, relations and values because of the easier, basic, common or easily understandable but highly inefficient solutions

Exercise 1.5: Inefficiently Finding Projects Active During a Specific Period

```

1 SELECT project_name, start_date, COALESCE(actual_end_date,
2     planned_end_date) AS relevant_end_date
3 FROM projects
4 WHERE
5     (start_date BETWEEN '2023-01-01' AND '2023-03-31') OR
6     (COALESCE(actual_end_date, planned_end_date) BETWEEN '2023-01-01' AND
7     '2023-03-31') OR
8     (start_date < '2023-01-01' AND COALESCE(actual_end_date,
9     planned_end_date) > '2023-03-31');

```

Listing 6: Solution for Exercise 1.5 - Inefficient Method

```

1 SELECT project_name, start_date, COALESCE(actual_end_date,
      planned_end_date) AS relevant_end_date
2 FROM projects
3 WHERE (start_date, COALESCE(actual_end_date, planned_end_date, DATE '
      9999-12-31'))
4         OVERLAPS (DATE '2023-01-01', DATE '2023-03-31');
5 -- Note: planned_end_date is NOT NULL in projects table.
6 -- If actual_end_date is NULL, planned_end_date is used.
7 -- If actual_end_date is present, it's used.
8 -- (start_date, COALESCE(actual_end_date, planned_end_date)) OVERLAPS (
      DATE '2023-01-01', DATE '2023-03-31');
9 -- The above is also correct given planned_end_date is NOT NULL.
10 -- The DATE '9999-12-31' handles cases where both could be NULL if
      schema allowed.

```

Listing 7: Solution for Exercise 1.5 - Efficient Method (using OVERLAPS)

Comparison: The OVERLAPS version is more concise, directly expresses the intent, and is less prone to logical errors.

2 Cases (Complementary SQL) - Solutions

Concepts: Searched CASE expressions, CASE in ORDER BY, CASE in GROUP BY.

(i) Meaning, values, relations, advantages of unique usage

Exercise 2.1: Project Status Categorization

```
1 SELECT
2     project_name ,
3     start_date ,
4     planned_end_date ,
5     actual_end_date ,
6     CASE
7         WHEN start_date > DATE '2024-01-15' THEN 'Upcoming'
8         WHEN start_date <= DATE '2024-01-15' AND (actual_end_date IS
9 NULL OR actual_end_date > DATE '2024-01-15') THEN 'Ongoing'
10        WHEN actual_end_date IS NOT NULL AND actual_end_date <=
11planned_end_date THEN 'Completed Early/On-Time'
12        WHEN actual_end_date IS NOT NULL AND actual_end_date >
13planned_end_date THEN 'Completed Late'
14        ELSE 'Status Unknown' -- e.g. completed project but
15        actual_end_date is null (bad data)
16    END AS project_status
17 FROM
18     projects;
```

Listing 8: Solution for Exercise 2.1

Exercise 2.2: Sorting Employees by Custom Priority

```
1 SELECT
2     emp_name ,
3     salary ,
4     hire_date ,
5     CASE
6         WHEN emp_id IN (SELECT DISTINCT manager_id FROM employees WHERE
7 manager_id IS NOT NULL) OR manager_id IS NULL THEN 'Manager/Top
8 Level'
9         ELSE 'Non-Manager'
10    END AS employee_category
11 FROM
12     employees
13 ORDER BY
14     CASE
15         WHEN emp_id IN (SELECT DISTINCT manager_id FROM employees WHERE
16 manager_id IS NOT NULL) OR manager_id IS NULL THEN 1 -- Managers
17 first
18         ELSE 2 -- Non-managers second
19     END ASC ,
20     CASE
21         WHEN emp_id IN (SELECT DISTINCT manager_id FROM employees WHERE
22 manager_id IS NOT NULL) OR manager_id IS NULL THEN salary
23         ELSE NULL
24     END DESC NULLS LAST ,
25     CASE
26         WHEN NOT (emp_id IN (SELECT DISTINCT manager_id FROM employees
27 WHERE manager_id IS NOT NULL) OR manager_id IS NULL) THEN hire_date
28         ELSE NULL
```

```
23 END ASC NULLS LAST;
```

Listing 9: Solution for Exercise 2.2

Exercise 2.3: Grouping Projects by Budget Ranges

```
1 SELECT
2     CASE
3         WHEN budget IS NULL THEN 'Undefined'
4         WHEN budget <= 50000 THEN 'Low'
5         WHEN budget > 50000 AND budget <= 150000 THEN 'Medium'
6         WHEN budget > 150000 THEN 'High'
7         ELSE 'Other' -- Should not happen with these conditions
8     END AS budget_category,
9     COUNT(project_id) AS number_of_projects,
10    SUM(budget) AS total_budget_in_category
11 FROM
12     projects
13 GROUP BY
14     budget_category -- Can group by alias in PostgreSQL
15 ORDER BY
16     MIN(budget) NULLS FIRST; -- Order categories by their typical
    budget size
```

Listing 10: Solution for Exercise 2.3

(ii) Disadvantages of all its technical concepts

Exercise 2.4: Overly Nested CASE Expressions for Readability

```
1 SELECT
2     emp_name,
3     salary,
4     dept_id,
5     performance_rating,
6     CASE
7         WHEN salary > 100000 THEN
8             'High Earner' ||
9             CASE
10                WHEN dept_id = 2 THEN
11                    ', Key Engineer' ||
12                    CASE WHEN performance_rating = 5 THEN ', Top
13                    Performer' ELSE '' END
14                ELSE ''
15            END ||
16            CASE -- Example: add something if not Eng but still high
17                earner and rating 5
18                WHEN dept_id != 2 AND performance_rating = 5 THEN ',
19                Non-Eng Top Performer'
20                ELSE ''
21            END
22        ELSE 'Standard Profile'
23    END AS profile_string
24 FROM
25     employees;
```

Listing 11: Solution for Exercise 2.4

Discussion of Disadvantage: Deeply nested CASE statements or very long flat ones can significantly reduce query readability and maintainability. Debugging becomes harder. Alternatives might involve CTEs or user-defined functions for very complex, reused logic.

Exercise 2.5: CASE in GROUP BY Causing Performance Issues with Non-SARGable Conditions

```
1 SELECT
2     CASE
3         WHEN email LIKE '%@example.com' THEN 'Internal Email'
4         ELSE 'External Email'
5     END AS email_category,
6     COUNT(*) AS number_of_employees
7 FROM
8     employees
9 GROUP BY
10    email_category; -- Can group by alias in PostgreSQL
```

Listing 12: Solution for Exercise 2.5

Discussion of Disadvantage: If GROUP BY CASE WHEN email LIKE '%@example.com' ... END is used:

1. **Index Inusability:** A standard B-tree index on email is unlikely to be used efficiently for a LIKE '%suffix' pattern. The database would likely scan all emails.
2. **Computational Cost:** Evaluating the CASE for every row before grouping adds overhead.

For better performance, functional indexes or specific text search index types might be needed. The disadvantage is potential poor performance on large tables without such optimizations.

(iii) Practice entirely cases where people in general does not use these approaches losing their advantages, relations and values because of the easier, basic, common or easily understandable but highly inefficient solutions

Exercise 2.6: Multiple UNION ALL Queries vs. CASE in GROUP BY for Segmented Counts

```
1 SELECT 'Engineering' AS department_segment, COUNT(*) AS employee_count
   FROM employees WHERE dept_id = 2
2 UNION ALL
3 SELECT 'Sales' AS department_segment, COUNT(*) AS employee_count FROM
   employees WHERE dept_id = 3
4 UNION ALL
5 SELECT 'Other Departments' AS department_segment, COUNT(*) AS
   employee_count FROM employees WHERE dept_id NOT IN (2, 3) OR dept_id
   IS NULL;
```

Listing 13: Solution for Exercise 2.6 - Inefficient Method (UNION ALL)

```
1 SELECT
2     CASE
3         WHEN dept_id = 2 THEN 'Engineering'
4         WHEN dept_id = 3 THEN 'Sales'
```

```

5         ELSE 'Other Departments'
6     END AS department_segment ,
7     COUNT(*) AS employee_count
8 FROM
9     employees
10 GROUP BY
11     department_segment; -- Can group by alias in PostgreSQL

```

Listing 14: Solution for Exercise 2.6 - Efficient Method (CASE in GROUP BY)

Comparison: The CASE in GROUP BY approach requires only one pass over the table, generally more efficient and concise than multiple UNION ALL operations.

3 Null Space (Complementary SQL) - Solutions

Concepts: NULLIF, NULL handling in aggregations, NULL handling in sorting.

(i) Meaning, values, relations, advantages of unique usage

Exercise 3.1: Safe Bonus Calculation Using NULLIF

```
1 SELECT
2     emp_name ,
3     salary ,
4     performance_rating ,
5     (salary * 0.10) / NULLIF(performance_rating, 1) AS
6     bonus_avoiding_rating_1
7 FROM
8     employees
9 WHERE
10    salary IS NOT NULL AND performance_rating IS NOT NULL;
```

Listing 15: Solution for Exercise 3.1

Exercise 3.2: Average Billing Rate Excluding Internal/Non-Billable Roles

```
1 SELECT
2     AVG(billing_rate) AS average_actual_billing_rate ,
3     COUNT(emp_project_id) AS total_assignments ,
4     COUNT(billing_rate) AS assignments_with_billing_rate
5 FROM
6     employee_projects;
```

Listing 16: Solution for Exercise 3.2

Discussion: AVG(billing_rate) calculates the sum of all non-NULL billing_rate values divided by the count of non-NULL billing_rate values. NULL billing rates are completely ignored. This is advantageous because:

1. It prevents NULLs from skewing the average (e.g., if NULL was treated as 0).
2. It reflects the average of roles/assignments that *are* billable.

If all billing_rates were NULL, AVG() would return NULL.

Exercise 3.3: Listing Projects by Actual End Date, Undefined Last

```
1 SELECT
2     project_name ,
3     actual_end_date
4 FROM
5     projects
6 ORDER BY
7     actual_end_date ASC NULLS LAST;
```

Listing 17: Solution for Exercise 3.3

(ii) Disadvantages of all its technical concepts

Exercise 3.4: NULLIF with Unintended Type Coercion or Comparison Issues

Discussion of Disadvantage & Example: If performance_rating is VARCHAR and we use NULLIF(performance_rating, 0):

1. **Type Coercion Error:** Many databases will try to convert either the string `performance_rating` to an integer, or the integer 0 to a string. If `performance_rating` contains non-numeric strings (e.g., 'Good', 'N/A'), converting it to an integer will fail.
2. **Unexpected Behavior (String Comparison):** If the integer 0 is converted to string '0', then `NULLIF('N/A', '0')` would not nullify 'N/A'.
3. **Portability:** Implicit type conversion rules can vary.

It's best practice to ensure `expr1` and `expr2` in `NULLIF` are of compatible types, or to use explicit `CAST`.

```

1 -- Hypothetical table and query illustrating the issue
2 CREATE TABLE IF NOT EXISTS reviews_temp (id INT, rating_text VARCHAR
  (10));
3 INSERT INTO reviews_temp VALUES (1, '5'), (2, 'N/A'), (3, '0');
4
5 -- This query in PostgreSQL will cast 0 to TEXT '0' for comparison
6 SELECT id, rating_text, NULLIF(rating_text, 0::TEXT) FROM reviews_temp;
7 -- Output will NULLIF '0', but not 'N/A'
8
9 -- If we wanted to NULLIF 'N/A':
10 SELECT id, rating_text, NULLIF(rating_text, 'N/A') FROM reviews_temp;
11
12 DROP TABLE IF EXISTS reviews_temp; -- Cleanup

```

Listing 18: Solution for Exercise 3.4 - Illustrative Example

The disadvantage is that carelessness with types in `NULLIF` can lead to runtime errors or logically incorrect nullifications.

Exercise 3.5: Aggregates over Mostly NULL Data Yielding Misleading Results

```

1 -- For HR dept (Frankenstein Monster: rating 3, Grace Hopper: rating
  NULL)
2 SELECT
3     d.dept_name,
4     AVG(e.performance_rating) AS avg_rating_ignoring_nulls,
5     COUNT(e.emp_id) AS total_employees,
6     COUNT(e.performance_rating) AS rated_employees
7 FROM
8     departments d
9 JOIN
10    employees e ON d.dept_id = e.dept_id
11 WHERE
12    d.dept_name = 'Human Resources'
13 GROUP BY
14    d.dept_name;

```

Listing 19: Solution for Exercise 3.5

Discussion: The "disadvantage" is potential for misinterpretation. For 'Human Resources', `avg_rating_ignoring_nulls` will be 3 (from Frankenstein Monster), while Grace Hopper's NULL rating is ignored. Reporting "Average rating of 3" alone can be misleading if the context (1 out of 2 employees rated) isn't provided. This calls for careful reporting, possibly including counts of total vs. contributing populations.

(iii) Practice entirely cases where people in general does not use these approaches losing their advantages, relations and values because of the easier, basic, common or easily understandable but highly inefficient solutions

Exercise 3.6: Using CASE WHEN expr = val THEN NULL ELSE expr END instead of NULLIF(expr, val)

```
1 SELECT
2     leave_id,
3     reason AS original_reason,
4     CASE
5         WHEN reason = '' THEN NULL
6         ELSE reason
7     END AS reason_via_case
8 FROM
9     leave_requests
10 WHERE emp_id = (SELECT emp_id from employees WHERE email='kevin@example
    .com');
```

Listing 20: Solution for Exercise 3.6 - Verbose CASE method

```
1 SELECT
2     leave_id,
3     reason AS original_reason,
4     NULLIF(reason, '') AS reason_via_nullif
5 FROM
6     leave_requests
7 WHERE emp_id = (SELECT emp_id from employees WHERE email='kevin@example
    .com');
```

Listing 21: Solution for Exercise 3.6 - Concise NULLIF method

Comparison: NULLIF(reason, '') is more direct and less verbose than the CASE expression for this specific task.

4 Hardcore Problem - Solution

Exercise 4.1: Comprehensive Departmental Project Health and Employee Engagement Report

```
1 WITH ReportContext AS (  
2     SELECT DATE '2024-01-15' AS current_report_date  
3 ),  
4 ActiveEmployees AS (  
5     SELECT  
6         e.emp_id, e.emp_name, e.dept_id, e.hire_date, e.  
7         performance_rating, e.manager_id,  
8         (rc.current_report_date - e.hire_date) / 365.25 AS  
9         tenure_years_raw  
10    FROM employees e, ReportContext rc  
11    WHERE (e.termination_date IS NULL OR e.termination_date > rc.  
12           current_report_date)  
13 ),  
14 ActiveManagersByDept AS (  
15     SELECT DISTINCT ae.dept_id, ae.emp_id AS manager_emp_id  
16    FROM ActiveEmployees ae  
17    WHERE ae.emp_id IN (SELECT DISTINCT manager_id FROM employees WHERE  
18           manager_id IS NOT NULL)  
19           OR ae.manager_id IS NULL  
20 ),  
21 ActiveProjects AS (  
22     SELECT p.project_id, p.project_name, p.start_date, p.  
23     planned_end_date, p.lead_emp_id  
24    FROM projects p, ReportContext rc  
25    WHERE (p.actual_end_date IS NULL OR p.actual_end_date > rc.  
26           current_report_date)  
27 ),  
28 DepartmentProjectAnalysis AS (  
29     SELECT  
30         d.dept_id,  
31         EXISTS (  
32             SELECT 1  
33             FROM ActiveProjects ap1  
34             JOIN ActiveManagersByDept amd1 ON ap1.lead_emp_id = amd1.  
35             manager_emp_id AND amd1.dept_id = d.dept_id  
36             JOIN ActiveProjects ap2 ON ap1.project_id < ap2.project_id  
37             JOIN ActiveManagersByDept amd2 ON ap2.lead_emp_id = amd2.  
38             manager_emp_id AND amd2.dept_id = d.dept_id  
39             WHERE (ap1.start_date, ap1.planned_end_date) OVERLAPS (ap2.  
40                    start_date, ap2.planned_end_date)  
41         ) AS has_overlap_risk,  
42         COUNT(DISTINCT ap.project_id) FILTER (  
43             WHERE ap.planned_end_date BETWEEN rc.current_report_date  
44             AND (rc.current_report_date + INTERVAL '30 days')  
45             AND EXISTS (SELECT 1 FROM ActiveManagersByDept amd WHERE  
46                    ap.lead_emp_id = amd.manager_emp_id AND amd.dept_id = d.dept_id)  
47         ) AS critical_deadline_projects_count  
48    FROM departments d  
49    LEFT JOIN ActiveProjects ap ON TRUE -- To allow access to rc and  
50    aggregate over d  
51    CROSS JOIN ReportContext rc -- Ensure rc.current_report_date is  
52    available
```



```

40     GROUP BY d.dept_id, rc.current_report_date -- Group by rc.
         current_report_date for correlated subquery access
41 ),
42 DepartmentAggregates AS (
43     SELECT
44         d.dept_id,
45         d.dept_name,
46         COUNT(DISTINCT ae.emp_id) AS total_active_employees,
47         AVG(ae.tenure_years_raw) AS avg_employee_tenure_years_raw,
48         AVG(COALESCE(ae.performance_rating, 2)) AS
avg_rating_adjusted_raw,
49         COUNT(DISTINCT CASE
50             WHEN lr.status = 'Approved' AND rc.
current_report_date BETWEEN lr.leave_start_date AND lr.
leave_end_date
51                 THEN ae.emp_id
52                 ELSE NULL
53             END
54         ) AS distinct_employees_on_leave
55     FROM departments d
56     LEFT JOIN ActiveEmployees ae ON d.dept_id = ae.dept_id
57     LEFT JOIN leave_requests lr ON ae.emp_id = lr.emp_id
58     CROSS JOIN ReportContext rc
59     GROUP BY d.dept_id, d.dept_name
60 )
61 SELECT
62     da.dept_name,
63     da.total_active_employees,
64     ROUND(da.avg_employee_tenure_years_raw::numeric, 2) AS
avg_employee_tenure_years,
65     CASE
66         WHEN dpa.has_overlap_risk THEN 'High Overlap Risk'
67         WHEN dpa.critical_deadline_projects_count > 1 THEN 'Multiple
Critical Deadlines'
68         ELSE 'Normal Load'
69     END AS projects_info,
70     ROUND(COALESCE(da.avg_rating_adjusted_raw, 0)::numeric, 2) AS
avg_rating_adjusted,
71     ROUND(
72         (da.distinct_employees_on_leave * 100.0) / NULLIF(da.
total_active_employees, 0),
73         2) AS employees_on_leave_percentage
74 FROM DepartmentAggregates da
75 JOIN DepartmentProjectAnalysis dpa ON da.dept_id = dpa.dept_id
76 ORDER BY
77     CASE
78         WHEN dpa.has_overlap_risk THEN 1
79         WHEN dpa.critical_deadline_projects_count > 1 THEN 2
80         ELSE 3
81     END ASC,
82     da.dept_name ASC
83 LIMIT 5;

```

Listing 22: Solution for Hardcore Exercise 4.1