

Set Operations and Subqueries

Complementary SQL: Solutions

May 15, 2025

Contents

1	Set Operations - Solutions	7
1.1	Type (i): Meaning, values, relations, advantages	7
1.1.1	SO-1.1: Unified List of All Current and Potential Engineering Staff	7
1.1.2	SO-1.2: Log of All Employee-Project Assignments and Recent Terminations	7
1.1.3	SO-1.3: Employees in Sales and Marketing Departments	7
1.1.4	SO-1.4: Active Engineers Not Assigned to 'Omega Security Update' Project	8
1.2	Type (ii): Disadvantages	8
1.2.1	SO-2.1: Mismatched Column Data Types in Union	8
1.2.2	SO-2.2: Performance of UNION vs UNION ALL with Large Datasets (Conceptual)	8
1.3	Type (iii): Inefficient alternatives	9
1.3.1	SO-3.1: Simulating EXCEPT with NOT IN (and its NULL pitfall)	9
1.3.2	SO-3.2: Simulating INTERSECT with Multiple Joins/WHERE conditions	9
1.4	Type (iv): Hardcore problem	10
1.4.1	SO-4.1: Consolidated List of High-Value Personnel Not On Leave, Ranked	10
2	Subqueries - Solutions	12
2.1	Type (i): Meaning, values, relations, advantages	12
2.1.1	SQ-1.1: Employees Earning More Than Average (Scalar Subquery)	12
2.1.2	SQ-1.2: Departments with Higher Than Average Project Budgets (Subquery in FROM)	12
2.1.3	SQ-1.3: Employee's Project Count (Subquery in SELECT - Correlated)	12
2.1.4	SQ-1.4: Employees in Departments Located in 'New York' (Subquery in WHERE with IN)	12
2.1.5	SQ-1.5: Departments with at Least One Project (Subquery in WHERE with EXISTS)	13
2.1.6	SQ-1.6: Employees Earning More Than Any Sales Intern (Subquery in WHERE with ANY/SOME)	13
2.1.7	SQ-1.7: Employees Earning More Than All Sales Interns (Subquery in WHERE with ALL)	13

2.2	Type (ii): Disadvantages	13
2.2.1	SQ-2.1: Performance of Correlated Subquery in SELECT	13
2.2.2	SQ-2.2: Scalar Subquery Returning Multiple Rows (Error Scenario)	14
2.3	Type (iii): Inefficient alternatives	14
2.3.1	SQ-3.1: Finding Max Salary Without Scalar Subquery (Inefficient Application Logic)	14
2.3.2	SQ-3.2: Checking Existence of Sales by Engineers (Inefficient: Fetch- ing All Sales Data)	14
2.4	Type (iv): Hardcore problem	15
2.4.1	SQ-4.1: Departmental High Earners Analysis	15

Global Dataset for PostgreSQL

The following SQL code creates and populates the necessary tables for the exercises. This dataset should be loaded into your PostgreSQL environment before running the solutions.

```
1  -- Drop tables if they exist to ensure a clean setup
2  DROP TABLE IF EXISTS EmployeeProjects CASCADE;
3  DROP TABLE IF EXISTS Sales CASCADE;
4  DROP TABLE IF EXISTS Products CASCADE;
5  DROP TABLE IF EXISTS OldEmployees CASCADE;
6  DROP TABLE IF EXISTS CandidateEmployees CASCADE;
7  DROP TABLE IF EXISTS OnLeaveEmployees CASCADE;
8  DROP TABLE IF EXISTS Employees CASCADE;
9  DROP TABLE IF EXISTS Departments CASCADE;
10 DROP TABLE IF EXISTS Projects CASCADE;
11
12 -- Create Tables
13 CREATE TABLE Departments (
14     departmentId INT PRIMARY KEY,
15     departmentName VARCHAR(100) NOT NULL UNIQUE,
16     locationCity VARCHAR(50)
17 );
18
19 CREATE TABLE Employees (
20     employeeId INT PRIMARY KEY,
21     firstName VARCHAR(50) NOT NULL,
22     lastName VARCHAR(50) NOT NULL,
23     email VARCHAR(100) UNIQUE,
24     phoneNumber VARCHAR(20),
25     hireDate DATE NOT NULL,
26     jobId VARCHAR(20),
27     salary NUMERIC(10, 2) NOT NULL CHECK (salary > 0),
28     commissionPct NUMERIC(4, 2) CHECK (commissionPct >= 0 AND commissionPct <= 1),
29     managerId INT REFERENCES Employees(employeeId),
30     departmentId INT REFERENCES Departments(departmentId)
31 );
32
33 CREATE TABLE Projects (
34     projectId INT PRIMARY KEY,
35     projectName VARCHAR(100) NOT NULL UNIQUE,
36     departmentId INT REFERENCES Departments(departmentId),
37     startDate DATE,
38     endDate DATE,
39     budget NUMERIC(12, 2) CHECK (budget >= 0)
40 );
41
42 CREATE TABLE EmployeeProjects (
43     employeeId INT REFERENCES Employees(employeeId),
44     projectId INT REFERENCES Projects(projectId),
45     assignedRole VARCHAR(50),
46     hoursWorked INT CHECK (hoursWorked >= 0),
47     PRIMARY KEY (employeeId, projectId)
48 );
49
50 CREATE TABLE OldEmployees (
51     employeeId INT PRIMARY KEY,
52     firstName VARCHAR(50),
53     lastName VARCHAR(50),
54     lastDepartmentId INT,
55     terminationDate DATE NOT NULL,
56     finalSalary NUMERIC(10, 2),
57     reasonForLeaving VARCHAR(255)
58 );
59
60 CREATE TABLE CandidateEmployees (
61     candidateId INT PRIMARY KEY,
62     firstName VARCHAR(50),
63     lastName VARCHAR(50),
64     appliedPosition VARCHAR(100),
65     expectedSalary NUMERIC(10, 2),
66     applicationDate DATE
67 );
```

```

68
69 CREATE TABLE OnLeaveEmployees (
70     employeeId INT PRIMARY KEY REFERENCES Employees(employeeId),
71     leaveStartDate DATE,
72     leaveEndDate DATE,
73     leaveReason VARCHAR(100)
74 );
75
76 CREATE TABLE Products (
77     productId INT PRIMARY KEY,
78     productName VARCHAR(100) NOT NULL,
79     productCategory VARCHAR(50),
80     unitPrice NUMERIC(10, 2) CHECK (unitPrice > 0)
81 );
82
83 CREATE TABLE Sales (
84     saleId INT PRIMARY KEY,
85     employeeId INT REFERENCES Employees(employeeId),
86     productId INT REFERENCES Products(productId),
87     saleDate TIMESTAMP NOT NULL,
88     quantitySold INT CHECK (quantitySold > 0),
89     saleAmount NUMERIC(10, 2)
90 );
91
92 -- Populate Tables
93 INSERT INTO Departments (departmentId, departmentName, locationCity) VALUES
94 (1, 'Human Resources', 'New York'),
95 (2, 'Engineering', 'San Francisco'),
96 (3, 'Sales', 'Chicago'),
97 (4, 'Marketing', 'New York'),
98 (5, 'Finance', 'London'),
99 (6, 'Research', 'San Francisco'),
100 (7, 'Customer Support', 'Austin');
101
102 INSERT INTO Employees (employeeId, firstName, lastName, email, phoneNumber, hireDate,
103     jobId, salary, commissionPct, managerId, departmentId) VALUES
104 (101, 'John', 'Smith', 'john.smith@example.com', '555-1234', '2018-06-01', 'HR_REP',
105     60000, NULL, NULL, 1),
106 (102, 'Alice', 'Johnson', 'alice.j@example.com', '555-5678', '2019-03-15', 'ENG_LEAD',
107     90000, NULL, NULL, 2),
108 (103, 'Bob', 'Williams', 'bob.w@example.com', '555-8765', '2019-07-20', 'SALES_MGR',
109     75000, 0.10, NULL, 3),
110 (104, 'Eva', 'Brown', 'eva.b@example.com', '555-4321', '2020-01-10', 'MKT_SPEC', 65000,
111     NULL, 101, 4),
112 (105, 'Charlie', 'Davis', 'charlie.d@example.com', '555-1122', '2018-11-05', 'ENG_SR',
113     85000, NULL, 102, 2),
114 (106, 'Diana', 'Miller', 'diana.m@example.com', '555-6543', '2021-05-25', 'SALES_REP',
115     55000, 0.05, 103, 3),
116 (107, 'Frank', 'Wilson', 'frank.w@example.com', '555-7890', '2022-08-01', 'ENG_JR',
117     70000, NULL, 102, 2),
118 (108, 'Grace', 'Moore', 'grace.m@example.com', '555-2109', '2019-09-01', 'FIN_ANALYST',
119     72000, NULL, NULL, 5),
120 (109, 'Henry', 'Taylor', 'henry.t@example.com', '555-1098', '2023-02-15', 'ENG_JR',
121     68000, NULL, 105, 2),
122 (110, 'Ivy', 'Anderson', 'ivy.a@example.com', '555-8076', '2020-11-30', 'MKT_MGR',
123     80000, NULL, 101, 4),
124 (111, 'Jack', 'Thomas', 'jack.t@example.com', '555-7654', '2017-07-14', 'RES_SCI',
125     95000, NULL, NULL, 6),
126 (112, 'Karen', 'Jackson', 'karen.j@example.com', '555-6547', '2021-10-01', 'HR_ASSIST',
127     50000, NULL, 101, 1),
128 (113, 'Leo', 'White', 'leo.w@example.com', '555-5438', '2023-05-20', 'SALES_REP', 58000,
129     0.06, 103, 3),
130 (114, 'Mia', 'Harris', 'mia.h@example.com', '555-4329', '2019-02-18', 'FIN_MGR', 92000,
131     NULL, 108, 5),
132 (115, 'Noah', 'Martin', 'noah.m@example.com', '555-3210', '2022-06-10', 'RES_ASSIST',
133     60000, NULL, 111, 6),
134 (116, 'Olivia', 'Garcia', 'olivia.g@example.com', '555-1987', '2018-09-01', 'ENG_SR',
135     88000, NULL, 102, 2),
136 (117, 'Paul', 'Martinez', 'paul.m@example.com', '555-8760', '2023-01-05', 'SALES_INTERN',
137     40000, 0.02, 106, 3),
138 (118, 'Quinn', 'Robinson', 'quinn.r@example.com', '555-7651', '2020-07-07', 'MKT_INTERN',
139     42000, NULL, 104, 4),

```

```

121 (119, 'Ruby', 'Clark', 'ruby.c@example.com', '555-6542', '2022-03-03', 'HR_SPEC', 62000,
    NULL, 101, 1),
122 (120, 'Sam', 'Rodriguez', 'sam.r@example.com', '555-5433', '2021-11-11', 'ENG_TECH',
    72000, NULL, 105, 2),
123 (121, 'Tom', 'Lee', 'tom.lee@example.com', '555-1122', '2023-08-15', 'FIN_ANALYST',
    73000, NULL, 114, 5),
124 (122, 'Ursula', 'Walker', 'ursula.w@example.com', '555-2233', '2019-01-20', 'RES_HEAD',
    120000, NULL, NULL, 6),
125 (123, 'Victor', 'Hall', 'victor.h@example.com', '555-3344', '2020-05-10', 'ENG_LEAD',
    95000, NULL, NULL, 2),
126 (124, 'Wendy', 'Allen', 'wendy.a@example.com', '555-4455', '2021-09-01', 'MKT_COORD',
    63000, NULL, 110, 4),
127 (125, 'Xavier', 'Young', 'xavier.y@example.com', '555-5566', '2022-12-12', 'SALES_LEAD',
    78000, 0.08, 103, 3),
128 (126, 'Yara', 'King', 'yara.k@example.com', '555-6677', '2018-04-04', 'HR_MGR', 85000,
    NULL, NULL, 1),
129 (127, 'Zack', 'Wright', 'zack.w@example.com', '555-7788', '2023-07-01', 'ENG_INTERN',
    45000, NULL, 107, 2),
130 (128, 'Laura', 'Palmer', 'laura.p@example.com', '555-1111', '2023-01-15', 'SUPPORT_REP',
    52000, NULL, NULL, 7),
131 (129, 'Dale', 'Cooper', 'dale.c@example.com', '555-2222', '2023-02-20', 'SUPPORT_LEAD',
    65000, NULL, NULL, 7);

132
133 UPDATE Employees SET managerId = 102 WHERE employeeId IN (105, 107, 116, 120);
134 UPDATE Employees SET managerId = 123 WHERE employeeId IN (109, 127);
135 UPDATE Employees SET managerId = 103 WHERE employeeId IN (106, 113, 117, 125);
136 UPDATE Employees SET managerId = 110 WHERE employeeId IN (104, 118, 124);
137 UPDATE Employees SET managerId = 101 WHERE employeeId IN (112, 119);
138 UPDATE Employees SET managerId = 126 WHERE employeeId = 101;
139 UPDATE Employees SET managerId = 111 WHERE employeeId = 115;
140 UPDATE Employees SET managerId = 122 WHERE employeeId = 111;
141 UPDATE Employees SET managerId = 114 WHERE employeeId IN (108, 121);
142 UPDATE Employees SET managerId = 129 WHERE employeeId = 128;
143
144 INSERT INTO Projects (projectId, projectName, departmentId, startDate, endDate, budget)
    VALUES
145 (1, 'Alpha Launch', 4, '2023-01-01', '2023-06-30', 150000.00),
146 (2, 'Beta Platform', 2, '2022-09-01', '2024-03-31', 500000.00),
147 (3, 'Gamma Sales Drive', 3, '2023-03-01', '2023-09-30', 80000.00),
148 (4, 'Delta HR System', 1, '2023-02-01', '2023-12-31', 120000.00),
149 (5, 'Epsilon Research', 6, '2022-05-01', '2024-05-30', 300000.00),
150 (6, 'Zeta Finance Tool', 5, '2023-07-01', '2024-06-30', 200000.00),
151 (7, 'Omega Security Update', 2, '2023-10-01', '2024-01-31', 250000.00),
152 (8, 'Sigma Marketing Campaign', 4, '2024-01-15', '2024-07-15', 180000.00),
153 (9, 'Kappa Efficiency Audit', 5, '2022-11-01', '2023-04-30', 75000.00),
154 (10, 'New Support Portal', 7, '2023-05-01', '2023-11-30', 90000.00);
155
156 INSERT INTO EmployeeProjects (employeeId, projectId, assignedRole, hoursWorked) VALUES
157 (102, 2, 'Project Lead', 500), (105, 2, 'Senior Developer', 600), (107, 2, 'Junior
    Developer', 450), (116, 2, 'Senior Developer', 550), (120, 2, 'Technician', 400),
    (123, 7, 'Project Lead', 200),
158 (104, 1, 'Marketing Specialist', 300), (110, 1, 'Campaign Manager', 250), (124, 1, '
    Coordinator', 320),
159 (106, 3, 'Sales Representative', 400), (113, 3, 'Sales Representative', 380), (125, 3, '
    Lead Sales', 350),
160 (101, 4, 'HR Lead', 200), (112, 4, 'HR Assistant', 300), (119, 4, 'HR Specialist', 280),
161 (111, 5, 'Lead Scientist', 700), (115, 5, 'Research Assistant', 650), (122, 5, '
    Principal Investigator', 500),
162 (108, 6, 'Financial Analyst', 300), (114, 6, 'Finance Lead', 250), (121, 6, 'Analyst',
    320),
163 (105, 7, 'Security Consultant', 150), (107, 7, 'Developer', 180),
164 (108, 9, 'Auditor', 200), (114, 9, 'Audit Lead', 150),
165 (128, 10, 'Support Analyst', 350), (129, 10, 'Project Manager', 280);
166
167 INSERT INTO OldEmployees (employeeId, firstName, lastName, lastDepartmentId,
    terminationDate, finalSalary, reasonForLeaving) VALUES
168 (201, 'Gary', 'Oldman', 2, '2022-12-31', 82000, 'Retired'),
169 (202, 'Helen', 'Hunt', 3, '2023-03-15', 70000, 'New Opportunity'),
170 (203, 'Mike', 'Myers', 2, '2021-08-20', 90000, 'Relocation'),
171 (204, 'Olivia', 'Garcia', 2, '2023-11-01', 88000, 'New Opportunity');
172
173 INSERT INTO CandidateEmployees (candidateId, firstName, lastName, appliedPosition,
    expectedSalary, applicationDate) VALUES

```

```

174 (301, 'Peter', 'Pan', 'ENG_JR', 65000, '2023-10-01'),
175 (302, 'Wendy', 'Darling', 'MKT_SPEC', 68000, '2023-09-15'),
176 (303, 'John', 'Smith', 'HR_REP', 60000, '2023-11-01'),
177 (304, 'Alice', 'Wonder', 'ENG_LEAD', 92000, '2023-08-20'),
178 (305, 'Bruce', 'Wayne', 'FIN_ANALYST', 75000, '2023-11-05');
179
180 INSERT INTO OnLeaveEmployees (employeeId, leaveStartDate, leaveEndDate, leaveReason)
      VALUES
181 (104, '2023-11-01', '2024-02-01', 'Maternity Leave'),
182 (111, '2023-09-15', '2023-12-15', 'Sabbatical');
183
184 INSERT INTO Products (productId, productName, productCategory, unitPrice) VALUES
185 (1, 'AlphaWidget', 'Electronics', 49.99),
186 (2, 'BetaGear', 'Software', 199.00),
187 (3, 'GammaCore', 'Hardware', 120.50),
188 (4, 'DeltaService', 'Services', 75.00),
189 (5, 'EpsilonPlus', 'Electronics', 89.90);
190
191 INSERT INTO Sales (saleId, employeeId, productId, saleDate, quantitySold, saleAmount)
      VALUES
192 (1, 106, 1, '2023-04-10 10:30:00', 2, 99.98),
193 (2, 106, 3, '2023-04-12 14:00:00', 1, 120.50),
194 (3, 113, 2, '2023-05-05 11:15:00', 1, 199.00),
195 (4, 106, 1, '2023-05-20 16:45:00', 3, 149.97),
196 (5, 117, 4, '2023-06-01 09:00:00', 10, 750.00),
197 (6, 125, 5, '2023-07-10 12:30:00', 5, 449.50),
198 (7, 113, 1, '2023-07-15 15:00:00', 2, 99.98),
199 (8, 103, 2, '2023-08-01 10:00:00', 2, 398.00),
200 (9, 106, 3, '2023-08-18 13:20:00', 1, 120.50),
201 (10, 125, 2, '2023-09-05 17:00:00', 1, 199.00),
202 (11, 113, 5, '2023-11-10 09:30:00', 3, 269.70),
203 (12, 106, 4, '2023-11-15 11:45:00', 5, 375.00),
204 (13, 117, 1, '2023-11-20 14:15:00', 1, 49.99),
205 (14, 125, 3, '2023-12-01 10:00:00', 2, 241.00),
206 (15, 103, 5, '2023-12-05 16:30:00', 4, 359.60),
207 (16, 128, 2, '2023-06-15 10:00:00', 1, 199.00),
208 (17, 105, 3, '2023-07-20 11:00:00', 1, 120.50);

```

Listing 1: Global Dataset for Exercises

1 Set Operations - Solutions

1.1 Type (i): Meaning, values, relations, advantages

1.1.1 SO-1.1: Unified List of All Current and Potential Engineering Staff

```
1 SELECT E.firstName || ' ' || E.lastName AS fullName, 'Current Employee'
   AS sourceType
2 FROM Employees E
3 JOIN Departments D ON E.departmentId = D.departmentId
4 WHERE D.departmentName = 'Engineering'
5 UNION
6 SELECT C.firstName || ' ' || C.lastName AS fullName, 'Candidate' AS
   sourceType
7 FROM CandidateEmployees C
8 WHERE C.appliedPosition IN ('ENG_LEAD', 'ENG_JR')
9 ORDER BY fullName;
```

This exercise shows UNION's ability to merge results from different queries with compatible columns and automatically remove duplicates. The advantage is a clean, consolidated list without manual deduplication.

1.1.2 SO-1.2: Log of All Employee-Project Assignments and Recent Terminations

```
1 SELECT employeeId, CAST(projectId AS VARCHAR) AS activityIdentifier, '
   Project Assignment' AS activityType
2 FROM EmployeeProjects
3 UNION ALL
4 SELECT employeeId, CAST(terminationDate AS VARCHAR) AS
   activityIdentifier, 'Termination' AS activityType
5 FROM OldEmployees
6 WHERE EXTRACT(YEAR FROM terminationDate) = 2023
7 ORDER BY employeeId, activityType;
```

This exercise highlights UNION ALL. Its advantage is speed, as it doesn't perform duplicate checking, which is suitable when all records from all sources are desired, like in a log.

1.1.3 SO-1.3: Employees in Sales and Marketing Departments

```
1 -- Employees in 'Sales' for 'Gamma Sales Drive'
2 SELECT E.employeeId, E.firstName, E.lastName
3 FROM Employees E
4 JOIN EmployeeProjects EP ON E.employeeId = EP.employeeId
5 JOIN Projects P ON EP.projectId = P.projectId
6 WHERE P.projectName = 'Gamma Sales Drive'
7 INTERSECT
8 -- Employees in 'Marketing' department
9 SELECT E.employeeId, E.firstName, E.lastName
10 FROM Employees E
11 JOIN Departments D ON E.departmentId = D.departmentId
12 WHERE D.departmentName = 'Marketing'
13 ORDER BY lastName, firstName;
```

This exercise shows INTERSECT for finding records present in both result sets. Its advantage is clearly expressing the need for common elements.

1.1.4 SO-1.4: Active Engineers Not Assigned to 'Omega Security Update' Project

```
1 -- All Engineers
2 SELECT E.employeeId, E.firstName, E.lastName
3 FROM Employees E
4 JOIN Departments D ON E.departmentId = D.departmentId
5 WHERE D.departmentName = 'Engineering'
6 EXCEPT
7 -- Engineers on 'Omega Security Update'
8 SELECT E.employeeId, E.firstName, E.lastName
9 FROM Employees E
10 JOIN EmployeeProjects EP ON E.employeeId = EP.employeeId
11 JOIN Projects P ON EP.projectId = P.projectId
12 WHERE P.projectName = 'Omega Security Update'
13 ORDER BY lastName, firstName;
```

*This exercise uses **EXCEPT** to subtract one set of employees from another. Its advantage is a concise way to find elements unique to the first query's result set.*

1.2 Type (ii): Disadvantages

1.2.1 SO-2.1: Mismatched Column Data Types in Union

```
1 -- This query will likely cause an error or implicit cast with
  -- potential issues:
2 -- SELECT employeeId, hireDate FROM Employees
3 -- UNION
4 -- SELECT candidateId, expectedSalary FROM CandidateEmployees;
5
6 -- Corrected version (if the goal was to show IDs and a numeric value):
7 SELECT employeeId AS id, salary AS amount, 'Employee' as source FROM
  Employees
8 WHERE departmentId = 2 -- Example filter
9 UNION
10 SELECT candidateId AS id, expectedSalary AS amount, 'Candidate' as
  source FROM CandidateEmployees
11 WHERE appliedPosition LIKE 'ENG%';
```

*Disadvantage: Set operations require that corresponding columns in the **SELECT** lists have compatible data types. A direct **UNION** of **hireDate** (**DATE**) and **expectedSalary** (**NUMERIC**) would fail or lead to undesirable implicit conversions. The solution shows how to select compatible types (or cast explicitly if combining conceptually different but type-compatible data).*

1.2.2 SO-2.2: Performance of UNION vs UNION ALL with Large Datasets (Conceptual)

```
1 -- Potentially faster if duplicates are acceptable or known to be few
2 SELECT firstName FROM Employees
3 UNION ALL
4 SELECT firstName FROM CandidateEmployees;
5
6 -- Potentially slower due to deduplication, especially with large
  -- tables
```



```

7 SELECT firstName FROM Employees
8 UNION
9 SELECT firstName FROM CandidateEmployees;

```

Disadvantage: UNION performs a distinct sort/hash operation to remove duplicates, which can be resource-intensive (CPU, memory, I/O) on large datasets compared to UNION ALL which simply concatenates results. If duplicates are acceptable or if the queries are designed such that no duplicates are possible, UNION ALL is generally more performant.

1.3 Type (iii): Inefficient alternatives

1.3.1 SO-3.1: Simulating EXCEPT with NOT IN (and its NULL pitfall)

Inefficient/Potentially Risky (if subquery could return NULLs for the compared column):

```

1 SELECT E.employeeId, E.firstName, E.lastName
2 FROM Employees E
3 JOIN Departments D ON E.departmentId = D.departmentId
4 WHERE D.departmentName = 'Engineering'
5 AND E.employeeId NOT IN (SELECT OL.employeeId FROM OnLeaveEmployees
6 OL);

```

More Efficient/Robust (using EXCEPT):

```

1 SELECT E.employeeId, E.firstName, E.lastName
2 FROM Employees E
3 JOIN Departments D ON E.departmentId = D.departmentId
4 WHERE D.departmentName = 'Engineering'
5 EXCEPT
6 SELECT OL.employeeId, E.firstName, E.lastName
7 FROM OnLeaveEmployees OL
8 JOIN Employees E ON OL.employeeId = E.employeeId;

```

Explanation: The NOT IN approach is common but has a pitfall: if the subquery for NOT IN returns even one NULL value, the entire NOT IN condition will evaluate to UNKNOWN or FALSE for all rows, potentially yielding incorrect results. EXCEPT (or NOT EXISTS) is generally safer and often more efficient for "anti-join" patterns. Here, EXCEPT clearly expresses the set difference and handles NULLs predictably.

1.3.2 SO-3.2: Simulating INTERSECT with Multiple Joins/WHERE conditions

Using INTERSECT:

```

1 SELECT E.employeeId, E.firstName
2 FROM Employees E
3 JOIN Departments D ON E.departmentId = D.departmentId
4 WHERE D.departmentName = 'Engineering'
5 INTERSECT
6 SELECT E.employeeId, E.firstName
7 FROM Employees E
8 JOIN EmployeeProjects EP ON E.employeeId = EP.employeeId
9 JOIN Projects P ON EP.projectId = P.projectId
10 WHERE P.projectName = 'Beta Platform'
11 ORDER BY employeeId;

```

Alternative using JOINS:

```

1 SELECT DISTINCT E.employeeId, E.firstName
2 FROM Employees E
3 JOIN Departments D ON E.departmentId = D.departmentId
4 JOIN EmployeeProjects EP ON E.employeeId = EP.employeeId
5 JOIN Projects P ON EP.projectId = P.projectId
6 WHERE D.departmentName = 'Engineering' AND P.projectName = 'Beta
   Platform'
7 ORDER BY E.employeeId;

```

Explanation: In this case, the JOIN approach is quite common and often efficient. However, if the conditions for being in each "set" were much more complex and involved different tables or logic, INTERSECT can make the query more readable by clearly separating the criteria for each set. The "inefficiency" of avoiding INTERSECT isn't always about raw performance but can be about code clarity and maintainability when set logic is primary.

1.4 Type (iv): Hardcore problem

1.4.1 SO-4.1: Consolidated List of High-Value Personnel Not On Leave, Ranked

```

1 WITH ValuablePersonnel AS (
2     SELECT
3         E.employeeId,
4         E.firstName,
5         E.lastName,
6         E.salary AS amount,
7         'Current Employee' AS personnelType
8     FROM Employees E
9     WHERE E.salary > 70000
10    UNION ALL
11    SELECT
12        OE.employeeId,
13        OE.firstName,
14        OE.lastName,
15        OE.finalSalary AS amount,
16        'Former Employee' AS personnelType
17    FROM OldEmployees OE
18    WHERE OE.finalSalary > 75000 AND OE.reasonForLeaving IN ('New
   Opportunity', 'Retired')
19 ),
20 ActiveValuablePersonnel AS (
21     SELECT vp.employeeId, vp.firstName, vp.lastName, vp.amount, vp.
   personnelType
22     FROM ValuablePersonnel vp
23     EXCEPT
24     SELECT ole.employeeId, vp2.firstName, vp2.lastName, vp2.amount, vp2
   .personnelType
25     FROM OnLeaveEmployees ole
26     JOIN ValuablePersonnel vp2 ON ole.employeeId = vp2.employeeId
27 )
28 SELECT
29     avp.employeeId,
30     avp.firstName,
31     avp.lastName,
32     avp.amount,

```

```
33     avp.personnelType ,
34     RANK() OVER (PARTITION BY avp.personnelType ORDER BY avp.amount
35     DESC) as salaryRankByType
36 FROM ActiveValuablePersonnel avp
37 ORDER BY avp.personnelType , salaryRankByType , avp.lastName;
```

*This hardcore problem uses **UNION ALL** to combine two distinct sets of personnel, **EXCEPT** to filter out those on leave, Common Table Expressions (CTEs) for readability and modularity, and a **RANK()** window function to rank them as required. It also involves basic joins and filtering conditions (**WHERE**).*

2 Subqueries - Solutions

2.1 Type (i): Meaning, values, relations, advantages

2.1.1 SQ-1.1: Employees Earning More Than Average (Scalar Subquery)

```
1 SELECT employeeId, firstName, lastName, salary
2 FROM Employees
3 WHERE salary > (SELECT AVG(salary) FROM Employees)
4 ORDER BY salary DESC;
```

Advantage: Scalar subqueries allow dynamic comparison values. The average salary is calculated at runtime and used for filtering.

2.1.2 SQ-1.2: Departments with Higher Than Average Project Budgets (Subquery in FROM)

```
1 SELECT
2     D.departmentName,
3     DeptAvgBudgets.avgBudget
4 FROM Departments D
5 JOIN (
6     SELECT departmentId, AVG(budget) AS avgBudget
7     FROM Projects
8     WHERE departmentId IS NOT NULL
9     GROUP BY departmentId
10 ) AS DeptAvgBudgets ON D.departmentId = DeptAvgBudgets.departmentId
11 WHERE DeptAvgBudgets.avgBudget > (SELECT AVG(budget) FROM Projects)
12 ORDER BY DeptAvgBudgets.avgBudget DESC;
```

Advantage: Subqueries in the FROM clause allow for pre-aggregation or transformation of data into an intermediate result set that can then be joined or filtered.

2.1.3 SQ-1.3: Employee's Project Count (Subquery in SELECT - Correlated)

```
1 SELECT
2     E.employeeId,
3     E.firstName,
4     E.lastName,
5     (SELECT COUNT(EP.projectId)
6      FROM EmployeeProjects EP
7      WHERE EP.employeeId = E.employeeId) AS projectCount
8 FROM Employees E
9 ORDER BY projectCount DESC, E.lastName;
```

Advantage: Correlated subqueries in SELECT can compute a specific value for each row of the outer query, useful for fetching related summary data.

2.1.4 SQ-1.4: Employees in Departments Located in 'New York' (Subquery in WHERE with IN)

```
1 SELECT employeeId, firstName, lastName
2 FROM Employees
```

```

3 WHERE departmentId IN (SELECT departmentId FROM Departments WHERE
   locationCity = 'New York')
4 ORDER BY lastName, firstName;

```

Advantage: IN with a subquery is concise for checking membership in a dynamically generated list of values.

2.1.5 SQ-1.5: Departments with at Least One Project (Subquery in WHERE with EXISTS)

```

1 SELECT D.departmentId, D.departmentName
2 FROM Departments D
3 WHERE EXISTS (SELECT 1 FROM Projects P WHERE P.departmentId = D.
   departmentId)
4 ORDER BY D.departmentName;

```

Advantage: EXISTS is efficient for checking the existence of related rows. It stops processing the subquery as soon as a match is found.

2.1.6 SQ-1.6: Employees Earning More Than Any Sales Intern (Subquery in WHERE with ANY/SOME)

```

1 SELECT employeeId, firstName, lastName, salary
2 FROM Employees
3 WHERE salary > ANY (SELECT salary FROM Employees WHERE jobId = '
   SALES_INTERN')
4 AND jobId != 'SALES_INTERN' -- Exclude sales interns themselves if
   desired
5 ORDER BY salary;

```

Advantage: ANY (or SOME) allows comparison against a set of values. > ANY means greater than the minimum value in the set returned by the subquery.

2.1.7 SQ-1.7: Employees Earning More Than All Sales Interns (Subquery in WHERE with ALL)

```

1 SELECT employeeId, firstName, lastName, salary
2 FROM Employees
3 WHERE salary > ALL (SELECT salary FROM Employees WHERE jobId = '
   SALES_INTERN')
4 ORDER BY salary;

```

Advantage: ALL allows comparison against a set of values. > ALL means greater than the maximum value in the set returned by the subquery.

2.2 Type (ii): Disadvantages

2.2.1 SQ-2.1: Performance of Correlated Subquery in SELECT

Solution (Correlated Subquery):

```

1 SELECT
2     E.employeeId,
3     E.firstName,
4     E.departmentId,
5     (SELECT P.projectName

```

```

6      FROM Projects P
7      WHERE P.departmentId = E.departmentId -- Correlation
8      ORDER BY P.budget DESC
9      LIMIT 1) AS mostExpensiveDeptProject
10 FROM Employees E
11 WHERE E.departmentId IS NOT NULL;

```

*Disadvantage: Correlated subqueries, especially in the **SELECT** list or **WHERE** clause, execute repeatedly for each row of the outer query. If the subquery is complex or operates on large tables, this can lead to poor performance.*

2.2.2 SQ-2.2: Scalar Subquery Returning Multiple Rows (Error Scenario)

```

1 -- This will error if the subquery returns more than one row:
2 -- SELECT employeeId, firstName, salary
3 -- FROM Employees
4 -- WHERE salary = (SELECT DISTINCT salary FROM Employees E JOIN
5 --                 Departments D ON E.departmentId = D.departmentId WHERE D.
6 --                 departmentName = 'Sales');
5
6 -- Corrected version using IN:
7 SELECT E.employeeId, E.firstName, E.salary
8 FROM Employees E
9 WHERE E.salary IN (SELECT DISTINCT Em.salary FROM Employees Em JOIN
10                  Departments Dp ON Em.departmentId = Dp.departmentId WHERE Dp.
11                  departmentName = 'Sales');

```

Disadvantage: A scalar subquery (one used where a single value is expected) must return exactly one row and one column. If it returns more than one row, it will result in a runtime error.

2.3 Type (iii): Inefficient alternatives

2.3.1 SQ-3.1: Finding Max Salary Without Scalar Subquery (Inefficient Application Logic)

Inefficient (conceptual application-level two-step): 1. **SELECT MAX(salary) FROM Employees;** (Result e.g., 120000) 2. **SELECT employeeId, firstName, salary FROM Employees WHERE salary = 120000;**

Efficient (using a scalar subquery):

```

1 SELECT employeeId, firstName, lastName, salary
2 FROM Employees
3 WHERE salary = (SELECT MAX(salary) FROM Employees);

```

Explanation: Performing two separate queries from the application adds network latency and complexity. A scalar subquery allows the database to handle this in a single, optimized operation.

2.3.2 SQ-3.2: Checking Existence of Sales by Engineers (Inefficient: Fetching All Sales Data)

Inefficient (conceptual: fetching lots of data then checking): **SELECT * FROM Sales;** (and then process in application to join with Engineers)

Efficient (using **EXISTS** with a correlated subquery):

```

1 SELECT D.departmentName
2 FROM Departments D
3 WHERE D.departmentName = 'Engineering'
4     AND EXISTS (SELECT 1
5                 FROM Sales S
6                 JOIN Employees E ON S.employeeId = E.employeeId
7                 WHERE E.departmentId = D.departmentId);
8 -- To list the engineers:
9 SELECT E.employeeId, E.firstName
10 FROM Employees E
11 JOIN Departments D ON E.departmentId = D.departmentId
12 WHERE D.departmentName = 'Engineering'
13     AND EXISTS (SELECT 1 FROM Sales S WHERE S.employeeId = E.employeeId);

```

*Explanation: The inefficient alternative involves fetching potentially large amounts of sales data. The **EXISTS** subquery is highly efficient because the database can stop searching as soon as the first matching sale is found.*

2.4 Type (iv): Hardcore problem

2.4.1 SQ-4.1: Departmental High Earners Analysis

```

1 WITH OverallAvgSalaryPost2020 AS (
2     SELECT AVG(salary) AS globalAvgSalary
3     FROM Employees
4     WHERE hireDate >= '2020-01-01'
5 ),
6 DepartmentalStats AS (
7     SELECT
8         D.departmentId,
9         D.departmentName,
10        AVG(E.salary) AS avgDeptSalary,
11        SUM(P.budget) AS totalDeptBudget
12    FROM Departments D
13    JOIN Employees E ON D.departmentId = E.departmentId
14    LEFT JOIN Projects P ON D.departmentId = P.departmentId
15    GROUP BY D.departmentId, D.departmentName
16    HAVING AVG(E.salary) > (SELECT globalAvgSalary FROM
17        OverallAvgSalaryPost2020)
18 ),
19 RankedEmployeesInQualifiedDepts AS (
20     SELECT
21         E.employeeId,
22         E.firstName,
23         E.lastName,
24         E.salary,
25         E.departmentId,
26         DS.departmentName,
27         DS.avgDeptSalary,
28         DS.totalDeptBudget,
29         ROW_NUMBER() OVER (PARTITION BY E.departmentId ORDER BY E.
30             salary DESC) AS rn
31     FROM Employees E
32     JOIN DepartmentalStats DS ON E.departmentId = DS.departmentId
33 )
34 SELECT
35     RE.departmentName,

```

```

34     RE.firstName,
35     RE.lastName,
36     RE.salary,
37     (RE.salary - RE.avgDeptSalary) AS salaryDiffFromDeptAvg,
38     COALESCE(RE.totalDeptBudget, 0) AS departmentTotalProjectBudget
39 FROM RankedEmployeesInQualifiedDepts RE
40 WHERE RE.rn <= 2
41 ORDER BY RE.departmentName, RE.salary DESC;

```

This hardcore problem utilizes: Scalar Subquery, Subquery in FROM (via CTEs), Aggregators, Joins, Window Functions, Basic SQL, and Null Handling to solve a multi-step analytical problem.