

Common Table Expressions - CTEs

Advanced Query Techniques: Solutions

May 17, 2025

Contents

1	Category (i): Solutions for Meanings, Values, Relations, and Advantages	6
1.1	Exercise 1: Basic CTE for Readability	6
1.2	Exercise 2: CTE Referenced Multiple Times	6
1.3	Exercise 3: Nested CTEs	6
1.4	Exercise 4: Recursive CTE for Hierarchical Data	7
2	Category (ii): Solutions for Disadvantages	8
2.1	Exercise 1: Potential Performance Issue (Optimization Fence / Materialization)	8
2.2	Exercise 2: No Indexing on CTE Results	8
2.3	Exercise 3: CTE Scope Limitation	9
3	Category (iii): Solutions for Cases Avoiding Inefficient Basic Solutions	10
3.1	Exercise 1: Replacing Repeated Subqueries	10
3.2	Exercise 2: Simplifying Complex Joins and Filters	10
3.3	Exercise 3: Avoiding Temporary Tables for Single-Query Scope	11
3.4	Exercise 4: Step-by-Step Multi-Level Aggregations (No Window Functions)	11
4	Category (iv): Solution for Hardcore Combined Problem	13

Global Dataset for PostgreSQL

The following SQL code creates and populates the necessary tables for the exercises. Execute this script in your PostgreSQL environment before attempting the exercises.

```
1  -- Dataset for Category (i)
2  CREATE TABLE DepartmentsI (
3      departmentId INTEGER PRIMARY KEY,
4      departmentName VARCHAR(100) NOT NULL,
5      locationCity VARCHAR(50)
6  );
7
8  CREATE TABLE EmployeesI (
9      employeeId INTEGER PRIMARY KEY,
10     employeeName VARCHAR(100) NOT NULL,
11     departmentId INTEGER REFERENCES DepartmentsI(departmentId),
12     managerId INTEGER REFERENCES EmployeesI(employeeId), -- Self-reference for hierarchy
13     salary DECIMAL(10, 2) NOT NULL,
14     hireDate DATE NOT NULL
15 );
16
17 INSERT INTO DepartmentsI (departmentId, departmentName, locationCity) VALUES
18 (1, 'Technology', 'New York'),
19 (2, 'Human Resources', 'London'),
20 (3, 'Sales', 'Tokyo'),
21 (4, 'Marketing', 'Paris');
22
23 INSERT INTO EmployeesI (employeeId, employeeName, departmentId, managerId, salary,
24     hireDate) VALUES
25 (101, 'Alice Wonderland', 1, NULL, 120000.00, '2018-03-15'),
26 (102, 'Bob The Builder', 1, 101, 90000.00, '2019-07-01'),
27 (103, 'Charlie Brown', 1, 102, 80000.00, '2020-01-10'),
28 (104, 'Diana Prince', 2, NULL, 110000.00, '2017-05-20'),
29 (105, 'Eve Harrington', 2, 104, 75000.00, '2021-02-28'),
30 (106, 'Frankenstein Monster', 3, NULL, 130000.00, '2018-11-01'),
31 (107, 'Grace OMalley', 3, 106, 85000.00, '2019-05-15'),
32 (108, 'Henry Jekyll', 3, 106, 82000.00, '2022-08-20'),
33 (109, 'Ivy Pepper', 1, 101, 95000.00, '2020-06-01'),
34 (110, 'John Doe', NULL, 101, 60000.00, '2023-01-15');
35
36 -- Dataset for Category (ii)
37 CREATE TABLE ProductCategoriesII (
38     categoryId SERIAL PRIMARY KEY,
39     categoryName VARCHAR(50) UNIQUE NOT NULL
40 );
41
42 CREATE TABLE ProductsII (
43     productId SERIAL PRIMARY KEY,
44     productName VARCHAR(100) NOT NULL,
45     categoryId INTEGER REFERENCES ProductCategoriesII(categoryId),
46     basePrice DECIMAL(10,2)
47 );
48
49 CREATE TABLE SalesTransactionsII (
50     transactionId SERIAL PRIMARY KEY,
51     productId INTEGER REFERENCES ProductsII(productId),
52     saleDate TIMESTAMP NOT NULL,
53     quantitySold INTEGER NOT NULL,
54     discount DECIMAL(3,2) DEFAULT 0.00
55 );
56
57 INSERT INTO ProductCategoriesII (categoryName) VALUES ('Electronics'), ('Books'), ('Home
58     Goods');
59 INSERT INTO ProductsII (productName, categoryId, basePrice) VALUES
60 ('Laptop Pro', 1, 1200.00), ('Quantum Physics Primer', 2, 25.00), ('Smart LED Bulb', 3,
61     15.00),
62 ('Desktop Gamer', 1, 1800.00), ('History of Time', 2, 20.00), ('Robotic Vacuum', 3,
63     300.00);
64
65 DO $$
66 DECLARE
67     i INT;
```

```

64     pId INT;
65     sDate TIMESTAMP;
66     qty INT;
67 BEGIN
68     FOR i IN 1..10000 LOOP
69         pId := (MOD(i, 6)) + 1;
70         sDate := CURRENT_TIMESTAMP - (MOD(i,365) || ' days')::INTERVAL - (MOD(i,24) || '
            hours')::INTERVAL;
71         qty := (MOD(i, 5)) + 1;
72         INSERT INTO SalesTransactionsII (productId, saleDate, quantitySold, discount)
73         VALUES (pId, sDate, qty, CASE WHEN MOD(i,10) = 0 THEN 0.05 ELSE 0.00 END);
74     END LOOP;
75 END $$;
76
77 UPDATE SalesTransactionsII
78 SET saleDate = CURRENT_DATE - INTERVAL '1 month' + (MOD(transactionId, 30) || ' days')::
    INTERVAL
79 WHERE MOD(productId, 2) = 0; -- Update some products to have recent sales
80
81 -- Dataset for Category (iii)
82 CREATE TABLE CustomersIII (
83     customerId SERIAL PRIMARY KEY,
84     customerName VARCHAR(100) NOT NULL,
85     registrationDate DATE,
86     city VARCHAR(50)
87 );
88
89 CREATE TABLE ProductsMasterIII (
90     productId SERIAL PRIMARY KEY,
91     productName VARCHAR(100),
92     category VARCHAR(50)
93 );
94
95 CREATE TABLE OrdersIII (
96     orderId SERIAL PRIMARY KEY,
97     customerId INTEGER REFERENCES CustomersIII(customerId),
98     orderDate DATE,
99     shipmentRegion VARCHAR(50)
100 );
101
102 CREATE TABLE OrderItemsIII (
103     orderItemId SERIAL PRIMARY KEY,
104     orderId INTEGER REFERENCES OrdersIII(orderId),
105     productId INTEGER REFERENCES ProductsMasterIII(productId),
106     quantity INTEGER,
107     pricePerUnit DECIMAL(10,2)
108 );
109
110 INSERT INTO CustomersIII (customerName, registrationDate, city) VALUES
111 ('Global Corp', '2020-01-15', 'New York'), ('Local Biz', '2021-06-01', 'London'),
112 ('Alpha Inc', '2019-11-20', 'Tokyo'), ('Beta LLC', '2022-03-10', 'New York');
113
114 INSERT INTO ProductsMasterIII (productName, category) VALUES
115 ('Widget A', 'Gadgets'), ('Widget B', 'Gizmos'), ('Service C', 'Services'), ('Tool D', '
    Tools');
116
117 INSERT INTO OrdersIII (customerId, orderDate, shipmentRegion) VALUES
118 (1, '2022-02-10', 'North America'), (2, '2022-03-15', 'Europe'),
119 (1, '2023-04-20', 'North America'), (3, '2023-05-05', 'Asia'),
120 (2, '2023-06-10', 'Europe'), (4, '2022-07-01', 'North America');
121
122 INSERT INTO OrderItemsIII (orderId, productId, quantity, pricePerUnit) VALUES
123 (1, 1, 10, 50.00), (1, 2, 5, 100.00), (2, 3, 1, 200.00),
124 (3, 1, 20, 45.00), (3, 4, 2, 150.00), (4, 2, 8, 95.00),
125 (5, 3, 2, 190.00), (6, 4, 3, 140.00);
126
127 -- Dataset for Category (iv)
128 CREATE TABLE DepartmentsIV (
129     departmentId SERIAL PRIMARY KEY,
130     departmentName VARCHAR(100) NOT NULL,
131     headEmployeeId INTEGER -- Nullable, to be cross-referenced with EmployeesIV
132 );
133

```

```

134 CREATE TABLE EmployeesIV (
135     employeeId SERIAL PRIMARY KEY,
136     employeeName VARCHAR(100) NOT NULL,
137     departmentId INTEGER REFERENCES DepartmentsIV(departmentId),
138     managerId INTEGER REFERENCES EmployeesIV(employeeId), -- For hierarchy
139     salary DECIMAL(10, 2) NOT NULL,
140     hireDate DATE NOT NULL
141 );
142
143 ALTER TABLE DepartmentsIV ADD CONSTRAINT fkHeadEmployee FOREIGN KEY (headEmployeeId)
    REFERENCES EmployeesIV(employeeId) DEFERRABLE INITIALLY DEFERRED;
144
145 INSERT INTO DepartmentsIV (departmentId, departmentName) VALUES
146 (1, 'Engineering'), (2, 'Product Management'), (3, 'Research & Development'), (4, '
    Operations');
147
148 INSERT INTO EmployeesIV (employeeId, employeeName, departmentId, managerId, salary,
    hireDate) VALUES
149 (1, 'Ava CEO', 1, NULL, 250000, '2015-01-01'),
150 (2, 'Brian Lead', 1, 1, 150000, '2018-06-01'),
151 (3, 'Chloe SeniorDev', 1, 2, 110000, '2020-03-15'),
152 (4, 'David JuniorDev', 1, 3, 75000, '2022-07-01'),
153 (5, 'Eli PMHead', 2, 1, 160000, '2017-09-01'),
154 (6, 'Fiona SeniorPM', 2, 5, 120000, '2020-11-01'),
155 (7, 'George PM', 2, 6, 85000, '2021-05-10'),
156 (8, 'Hannah RDHead', 3, 1, 170000, '2016-04-12'),
157 (9, 'Ian SeniorScientist', 3, 8, 130000, '2021-01-20'),
158 (10, 'Julia Scientist', 3, 9, 90000, '2022-08-01'),
159 (11, 'Kevin OpsLead', 4, 1, 140000, '2019-02-10'),
160 (12, 'Liam OpsSpecialist', 4, 11, 95000, '2021-10-05'),
161 (13, 'Mike AnotherDev', 1, 2, 105000, '2021-02-01');
162
163 UPDATE DepartmentsIV SET headEmployeeId = 2 WHERE departmentName = 'Engineering';
164 UPDATE DepartmentsIV SET headEmployeeId = 5 WHERE departmentName = 'Product Management';
165 UPDATE DepartmentsIV SET headEmployeeId = 8 WHERE departmentName = 'Research &
    Development';
166 UPDATE DepartmentsIV SET headEmployeeId = 11 WHERE departmentName = 'Operations';
167
168 CREATE TABLE ProjectsIV (
169     projectId SERIAL PRIMARY KEY,
170     projectName VARCHAR(150) NOT NULL,
171     startDate DATE,
172     endDate DATE,
173     budget DECIMAL(12, 2)
174 );
175
176 CREATE TABLE TasksIV (
177     taskId SERIAL PRIMARY KEY,
178     projectId INTEGER REFERENCES ProjectsIV(projectId),
179     taskName VARCHAR(200),
180     assignedToEmployeeId INTEGER REFERENCES EmployeesIV(employeeId),
181     estimatedHours INTEGER,
182     actualHours INTEGER,
183     status VARCHAR(20)
184 );
185
186 CREATE TABLE TimeLogsIV (
187     logId SERIAL PRIMARY KEY,
188     taskId INTEGER REFERENCES TasksIV(taskId),
189     employeeId INTEGER REFERENCES EmployeesIV(employeeId),
190     logDate DATE NOT NULL,
191     hoursWorked DECIMAL(5,2) NOT NULL,
192     notes TEXT
193 );
194
195 INSERT INTO ProjectsIV (projectName, startDate, endDate, budget) VALUES
196 ('Alpha Core System', '2022-01-01', '2023-12-31', 200000.00),
197 ('Beta Mobile App', '2023-03-01', '2024-02-28', 80000.00),
198 ('Gamma Research Initiative', '2021-06-15', '2023-05-30', 160000.00),
199 ('Delta Operations Upgrade', '2023-07-01', NULL, 120000.00);
200
201 INSERT INTO TasksIV (projectId, taskName, assignedToEmployeeId, estimatedHours,
    actualHours, status) VALUES

```

```

202 (1, 'Design Alpha Architecture', 3, 100, 90, 'Completed'),
203 (1, 'Develop Alpha Module 1', 3, 150, 160, 'In Progress'),
204 (2, 'Beta UI/UX Design', 6, 80, 70, 'Completed'),
205 (2, 'Beta Backend Dev', 7, 120, 50, 'In Progress'),
206 (3, 'Gamma Initial Research', 9, 200, 180, 'Completed'),
207 (3, 'Gamma Experiment Setup', 9, 100, 110, 'Overdue'),
208 (4, 'Delta Process Analysis', 12, 60, 40, 'In Progress'),
209 (1, 'Alpha Documentation', 13, 80, 0, 'Pending');
210 -- The next task inserted will have taskId = 9
211 INSERT INTO TasksIV (projectId, taskName, assignedToEmployeeId, estimatedHours,
    actualHours, status) VALUES
212 (2, 'Cross-project review for Alpha', 3, 20, 0, 'Pending');
213
214 INSERT INTO TimeLogsIV (logId, taskId, employeeId, logDate, hoursWorked, notes) VALUES
215 (DEFAULT, 1, 3, '2022-03-01', 8.0, 'Initial design'), (DEFAULT, 1, 3, '2022-03-02', 8.0,
    'Refinement'),
216 (DEFAULT, 2, 3, '2022-04-01', 8.0, 'Dev start'), (DEFAULT, 2, 3, '2022-04-02', 8.0, '
    Core logic'),
217 (DEFAULT, 3, 6, '2023-03-10', 7.0, 'UX flows'),
218 (DEFAULT, 5, 9, '2021-07-01', 6.0, 'Literature review'), (DEFAULT, 5, 9, '2021-07-02',
    8.0, 'Planning'),
219 (DEFAULT, 6, 9, '2021-09-01', 8.0, 'Setup phase 1'), (DEFAULT, 6, 9, '2021-09-02', 5.0,
    'Troubleshooting setup'),
220 (DEFAULT, 7, 12, '2023-07-15', 8.0, 'Mapping current state'),
221 (DEFAULT, 8, 13, '2022-05-01', 4.0, 'Doc outline');
222 -- logId values will be 1 to 11 after these inserts
223 INSERT INTO TimeLogsIV (logId, taskId, employeeId, logDate, hoursWorked, notes) VALUES
224 (DEFAULT, 2, 3, '2022-04-03', 8.0, 'Task 2 for emp 3'), -- emp 3 (Chloe) on task 2 (
    project 1), logId 12
225 (DEFAULT, 4, 7, '2023-08-01', 5.0, 'Task 4 for emp 7'), -- emp 7 (George) on task 4 (
    project 2), logId 13
226 (DEFAULT, 9, 3, '2023-09-01', 3.0, 'Time for task 9, project 2'); -- emp 3 works on task
    9 (project 2), logId 14

```

Listing 1: Global Dataset for Exercises

1 Category (i): Solutions for Meanings, Values, Relations, and Advantages

1.1 Exercise 1: Basic CTE for Readability

```
1 WITH TechDepartment AS (  
2     SELECT departmentId  
3     FROM DepartmentsI  
4     WHERE departmentName = 'Technology'  
5 )  
6 SELECT e.employeeName, e.salary, d.departmentName  
7 FROM EmployeesI e  
8 JOIN TechDepartment td ON e.departmentId = td.departmentId  
9 JOIN DepartmentsI d ON e.departmentId = d.departmentId  
10 WHERE e.salary > 90000.00;
```

Listing 2: Solution for Exercise 1.1

1.2 Exercise 2: CTE Referenced Multiple Times

```
1 WITH DepartmentAvgSalary AS (  
2     SELECT  
3         departmentId,  
4         AVG(salary) AS avgSalaryForDept  
5     FROM EmployeesI  
6     WHERE departmentId IS NOT NULL  
7     GROUP BY departmentId  
8 )  
9 SELECT  
10     e.employeeName,  
11     e.salary,  
12     d.departmentName,  
13     das.avgSalaryForDept  
14 FROM EmployeesI e  
15 JOIN DepartmentsI d ON e.departmentId = d.departmentId  
16 JOIN DepartmentAvgSalary das ON e.departmentId = das.departmentId  
17 WHERE e.salary > das.avgSalaryForDept  
18 ORDER BY d.departmentName, e.employeeName;
```

Listing 3: Solution for Exercise 1.2

1.3 Exercise 3: Nested CTEs

```
1 WITH RelevantDepartments AS (  
2     SELECT departmentId, departmentName, locationCity  
3     FROM DepartmentsI  
4     WHERE locationCity IN ('New York', 'London')  
5 ), FilteredEmployees AS (  
6     SELECT employeeName, salary, departmentId, hireDate  
7     FROM EmployeesI  
8     WHERE EXTRACT(YEAR FROM hireDate) > 2019  
9 )  
10 SELECT  
11     fe.employeeName,
```

```

12     fe.salary,
13     rd.departmentName,
14     rd.locationCity,
15     fe.hireDate
16 FROM FilteredEmployees fe
17 JOIN RelevantDepartments rd ON fe.departmentId = rd.departmentId
18 ORDER BY rd.departmentName, fe.employeeName;

```

Listing 4: Solution for Exercise 1.3

1.4 Exercise 4: Recursive CTE for Hierarchical Data

```

1 WITH RECURSIVE EmployeeHierarchy AS (
2     SELECT
3         employeeId,
4         employeeName,
5         managerId,
6         0 AS level
7     FROM EmployeesI
8     WHERE employeeId = 103 -- Starting employee (Charlie Brown)
9
10    UNION ALL
11
12    SELECT
13        e.employeeId,
14        e.employeeName,
15        e.managerId,
16        eh.level + 1
17    FROM EmployeesI e
18    JOIN EmployeeHierarchy eh ON e.employeeId = eh.managerId
19 )
20 SELECT employeeId, employeeName, managerId, level
21 FROM EmployeeHierarchy
22 ORDER BY level DESC;

```

Listing 5: Solution for Exercise 1.4

2 Category (ii): Solutions for Disadvantages

2.1 Exercise 1: Potential Performance Issue (Optimization Fence / Materialization)

```
1 WITH ProductRevenue AS (  
2     SELECT  
3         st.productId,  
4         SUM(p.basePrice * st.quantitySold * (1 - st.discount)) AS  
5         totalRevenue  
6     FROM SalesTransactionsII st  
7     JOIN ProductsII p ON st.productId = p.productId  
8     GROUP BY st.productId -- This computes revenue for ALL products  
9 )  
10 SELECT  
11     p.productName,  
12     pr.totalRevenue  
13 FROM ProductRevenue pr  
14 JOIN ProductsII p ON pr.productId = p.productId  
15 JOIN ProductCategoriesII pc ON p.categoryId = pc.categoryId  
16 WHERE pc.categoryName = 'Electronics';  
17 -- Note: The disadvantage is that the CTE 'ProductRevenue' might be  
18 -- fully computed  
19 -- before the 'WHERE pc.categoryName = 'Electronics'' filter is applied  
20 -- , especially  
21 -- if the CTE is complex or the optimizer chooses to materialize it.
```

Listing 6: Solution for Exercise 2.1

2.2 Exercise 2: No Indexing on CTE Results

```
1 -- The dataset setup ensures some sales are in the previous month.  
2 -- Example: If current month is February, this looks for January sales.  
3 WITH PreviousMonthSalesProducts AS (  
4     SELECT DISTINCT productId  
5     FROM SalesTransactionsII  
6     WHERE DATE_TRUNC('month', saleDate) = DATE_TRUNC('month',  
7         CURRENT_DATE - INTERVAL '1 month')  
8 )  
9 -- Simulating multiple uses of the intermediate "Previous Month Sales  
10 -- Products" concept:  
11 -- Use 1: List product names  
12 SELECT p.productName  
13 FROM ProductsII p  
14 WHERE p.productId IN (SELECT productId FROM PreviousMonthSalesProducts)  
15 UNION ALL  
16 -- Use 2: Count these products (conceptually another query part)  
17 SELECT CONCAT('Total Previous Month Sales Products: ', COUNT(productId)  
18     ::TEXT)  
19 FROM PreviousMonthSalesProducts;  
20 -- Note: The disadvantage is that 'PreviousMonthSalesProducts' logic is  
21 -- either re-run or its temporary,  
22 -- unindexed result is scanned. If this CTE were used many times in a  
23 -- very complex query,  
24 -- this could be inefficient compared to a (materialized and indexed)  
25 -- temporary table
```



```
20 -- in procedural contexts.
```

Listing 7: Solution for Exercise 2.2

2.3 Exercise 3: CTE Scope Limitation

```
1  -- Query 1: Calculate total sales for Books (This works)
2  WITH BookSales AS (
3      SELECT SUM(p.basePrice * st.quantitySold * (1 - st.discount)) AS
4          totalBookRevenue
5      FROM SalesTransactionsII st
6      JOIN ProductsII p ON st.productId = p.productId
7      JOIN ProductCategoriesII pc ON p.categoryId = pc.categoryId
8      WHERE pc.categoryName = 'Books'
9  )
10 SELECT totalBookRevenue FROM BookSales;
11
12 -- Query 2: Attempt to use BookSales CTE (this will fail, illustrating
13 -- scope)
14 -- SELECT totalBookRevenue * 0.1 AS tenPercentOfBookRevenue FROM
15 -- BookSales;
16 -- The above query would cause an error: "ERROR: relation "booksales"
17 -- does not exist"
18 -- This demonstrates the scope limitation.
19
20 -- To achieve the desired outcome, the CTE must be re-declared if used
21 -- in a separate statement:
22 WITH BookSales AS (
23     SELECT SUM(p.basePrice * st.quantitySold * (1 - st.discount)) AS
24         totalBookRevenue
25     FROM SalesTransactionsII st
26     JOIN ProductsII p ON st.productId = p.productId
27     JOIN ProductCategoriesII pc ON p.categoryId = pc.categoryId
28     WHERE pc.categoryName = 'Books'
29 )
30 SELECT totalBookRevenue * 0.1 AS approximateProfit FROM BookSales;
31 -- Note: The point is that 'BookSales' had to be defined again.
32 -- If it were a very complex CTE, this is redundant.
```

Listing 8: Solution for Exercise 2.3

3 Category (iii): Solutions for Cases Avoiding Inefficient Basic Solutions

3.1 Exercise 1: Replacing Repeated Subqueries

```
1 WITH Orders2022 AS (  
2     SELECT DISTINCT customerId  
3     FROM OrdersIII  
4     WHERE EXTRACT(YEAR FROM orderDate) = 2022  
5 ), Orders2023 AS (  
6     SELECT DISTINCT customerId  
7     FROM OrdersIII  
8     WHERE EXTRACT(YEAR FROM orderDate) = 2023  
9 )  
10 SELECT c.customerName, c.city  
11 FROM CustomersIII c  
12 WHERE c.customerId IN (SELECT customerId FROM Orders2022)  
13 AND c.customerId IN (SELECT customerId FROM Orders2023);
```

Listing 9: Solution for Exercise 3.1

Alternatively, a slightly different CTE structure could also work:

```
1 WITH CustomerYearlyOrders AS (  
2     SELECT customerId, EXTRACT(YEAR FROM orderDate) AS orderYear  
3     FROM OrdersIII  
4     WHERE EXTRACT(YEAR FROM orderDate) IN (2022, 2023)  
5     GROUP BY customerId, EXTRACT(YEAR FROM orderDate)  
6 )  
7 SELECT c.customerName, c.city  
8 FROM CustomersIII c  
9 WHERE EXISTS (SELECT 1 FROM CustomerYearlyOrders cyo WHERE cyo.  
10     customerId = c.customerId AND cyo.orderYear = 2022)  
11 AND EXISTS (SELECT 1 FROM CustomerYearlyOrders cyo WHERE cyo.  
12     customerId = c.customerId AND cyo.orderYear = 2023);
```

Listing 10: Alternative Solution for Exercise 3.1

3.2 Exercise 2: Simplifying Complex Joins and Filters

```
1 WITH OrderTotals AS (  
2     SELECT  
3         orderId,  
4         SUM(quantity * pricePerUnit) AS totalValue  
5     FROM OrderItemsIII  
6     GROUP BY orderId  
7 ), NorthAmericaHighValueOrders AS (  
8     SELECT o.orderId  
9     FROM OrdersIII o  
10    JOIN OrderTotals ot ON o.orderId = ot.orderId  
11    WHERE o.shipmentRegion = 'North America' AND ot.totalValue > 600.00  
12 )  
13 SELECT DISTINCT p.productName, p.category  
14 FROM ProductsMasterIII p  
15 JOIN OrderItemsIII oi ON p.productId = oi.productId  
16 WHERE oi.orderId IN (SELECT orderId FROM NorthAmericaHighValueOrders)
```

```
17 ORDER BY p.productName;
```

Listing 11: Solution for Exercise 3.2

3.3 Exercise 3: Avoiding Temporary Tables for Single-Query Scope

```
1 WITH OrderValues AS (
2     SELECT
3         o.orderId,
4         o.shipmentRegion,
5         SUM(oi.quantity * oi.pricePerUnit) AS orderTotal
6     FROM OrdersIII o
7     JOIN OrderItemsIII oi ON o.orderId = oi.orderId
8     GROUP BY o.orderId, o.shipmentRegion
9 ), RegionalAverageOrderValue AS (
10    SELECT
11        shipmentRegion,
12        AVG(orderTotal) AS avgRegionalValue
13    FROM OrderValues
14    GROUP BY shipmentRegion
15 ), OverallAverageOrderValue AS (
16    SELECT AVG(orderTotal) AS overallAvgValue
17    FROM OrderValues
18 )
19 SELECT
20     raov.shipmentRegion,
21     raov.avgRegionalValue
22 FROM RegionalAverageOrderValue raov, OverallAverageOrderValue oaov
23 WHERE raov.avgRegionalValue > oaov.overallAvgValue
24 ORDER BY raov.avgRegionalValue DESC;
```

Listing 12: Solution for Exercise 3.3

3.4 Exercise 4: Step-by-Step Multi-Level Aggregations (No Window Functions)

```
1 WITH MonthlyCategorySales AS (
2     SELECT
3         p.category,
4         TO_CHAR(o.orderDate, 'YYYY-MM') AS saleMonth,
5         SUM(oi.quantity) AS totalQuantity
6     FROM OrderItemsIII oi
7     JOIN ProductsMasterIII p ON oi.productId = p.productId
8     JOIN OrdersIII o ON oi.orderId = o.orderId
9     GROUP BY p.category, TO_CHAR(o.orderDate, 'YYYY-MM')
10 ),
11 MaxQuantityPerCategory AS (
12     SELECT
13         category,
14         MAX(totalQuantity) AS maxTotalQuantity
15     FROM MonthlyCategorySales
16     GROUP BY category
17 )
18 SELECT
```

```
19     mcs.category ,
20     mcs.saleMonth ,
21     mcs.totalQuantity
22 FROM MonthlyCategorySales mcs
23 JOIN MaxQuantityPerCategory mqpc ON mcs.category = mqpc.category AND
    mcs.totalQuantity = mqpc.maxTotalQuantity
24 ORDER BY mcs.category , mcs.saleMonth;
```

Listing 13: Solution for Exercise 3.4

4 Category (iv): Solution for Hardcore Combined Problem

```
1 WITH CriticalProjects AS (  
2     SELECT projectId  
3     FROM ProjectsIV  
4     WHERE budget > 150000.00  
5 ),  
6 EmployeeLoggedTimeOnCriticalTasks AS ( -- Renamed for clarity and  
    corrected logic  
7     SELECT DISTINCT tl.employeeId  
8     FROM TimeLogsIV tl  
9     JOIN TasksIV t ON tl.taskId = t.taskId  
10    WHERE t.projectId IN (SELECT projectId FROM CriticalProjects)  
11 ),  
12 SeniorEmployees AS (  
13     SELECT  
14         e.employeeId,  
15         e.employeeName,  
16         e.departmentId,  
17         e.salary,  
18         e.hireDate,  
19         d.departmentName  
20     FROM EmployeesIV e  
21     JOIN DepartmentsIV d ON e.departmentId = d.departmentId  
22     WHERE e.salary > 70000.00  
23         AND e.hireDate >= '2020-01-01'  
24         AND e.employeeId IN (SELECT employeeId FROM  
    EmployeeLoggedTimeOnCriticalTasks)  
25 ),  
26 DepartmentSeniorStats AS (  
27     SELECT  
28         s.departmentId,  
29         s.departmentName,  
30         SUM(s.salary) AS totalSeniorSalary,  
31         COUNT(s.employeeId) AS countSeniorEmployees  
32     FROM SeniorEmployees s  
33     GROUP BY s.departmentId, s.departmentName  
34     HAVING COUNT(s.employeeId) >= 2  
35 ),  
36 TopDepartments AS (  
37     SELECT  
38         departmentId,  
39         departmentName,  
40         totalSeniorSalary,  
41         countSeniorEmployees  
42     FROM DepartmentSeniorStats  
43     ORDER BY totalSeniorSalary DESC  
44     FETCH FIRST 2 ROWS ONLY  
45 ),  
46 EmployeeProjectCounts AS (  
47     SELECT  
48         tl.employeeId,  
49         e.departmentId, -- DepartmentId of the employee logging time  
50         COUNT(DISTINCT t.projectId) AS distinctProjectCount  
51     FROM TimeLogsIV tl  
52     JOIN TasksIV t ON tl.taskId = t.taskId
```

```

53 JOIN EmployeesIV e ON tl.employeeId = e.employeeId -- Join to get
54 employee's department
55 GROUP BY tl.employeeId, e.departmentId
56 )
57 SELECT
58     td.departmentName,
59     td.totalSeniorSalary,
60     td.countSeniorEmployees,
61     empDetails.employeeName AS mostActiveEmployeeName,
62     empDetails.distinctProjectCount AS mostActiveEmployeeProjectCount
63 FROM TopDepartments td
64 LEFT JOIN LATERAL (
65     SELECT
66         epc.employeeId,
67         eiv.employeeName,
68         epc.distinctProjectCount
69     FROM EmployeeProjectCounts epc
70     JOIN EmployeesIV eiv ON epc.employeeId = eiv.employeeId
71     WHERE epc.departmentId = td.departmentId -- Match employee's
72     department with the top department
73     ORDER BY epc.distinctProjectCount DESC, epc.employeeId ASC
74     LIMIT 1
75 ) empDetails ON TRUE
76 ORDER BY td.totalSeniorSalary DESC;

```

Listing 14: Solution for Hardcore Problem (Part 1: Top Departments Info)

```

1 WITH CriticalProjects AS ( -- Re-declare or ensure CTEs are available
2     if running separately
3     SELECT projectId
4     FROM ProjectsIV
5     WHERE budget > 150000.00
6 ),
7 EmployeeLoggedTimeOnCriticalTasks AS (
8     SELECT DISTINCT tl.employeeId
9     FROM TimeLogsIV tl
10    JOIN TasksIV t ON tl.taskId = t.taskId
11    WHERE t.projectId IN (SELECT projectId FROM CriticalProjects)
12 ),
13 SeniorEmployees AS (
14     SELECT
15         e.employeeId,
16         e.departmentId,
17         e.salary,
18         e.hireDate
19     FROM EmployeesIV e
20     WHERE e.salary > 70000.00
21     AND e.hireDate >= '2020-01-01'
22     AND e.employeeId IN (SELECT employeeId FROM
23         EmployeeLoggedTimeOnCriticalTasks)
24 ),
25 DepartmentSeniorStats AS (
26     SELECT
27         s.departmentId,
28         SUM(s.salary) AS totalSeniorSalary
29     FROM SeniorEmployees s
30     GROUP BY s.departmentId
31     HAVING COUNT(s.employeeId) >= 2
32 );

```

```

31 TopDepartmentForHierarchy AS ( -- Select only the #1 department
32     SELECT
33         ds.departmentId,
34         d.departmentName,
35         d.headEmployeeId
36     FROM DepartmentSeniorStats ds
37     JOIN DepartmentsIV d ON ds.departmentId = d.departmentId
38     ORDER BY ds.totalSeniorSalary DESC
39     FETCH FIRST 1 ROWS ONLY
40 ),
41 EmployeeHierarchy AS (
42     SELECT
43         e.employeeId,
44         e.employeeName,
45         e.managerId,
46         0 AS level
47     FROM EmployeesIV e
48     WHERE e.employeeId = (SELECT headEmployeeId FROM
49         TopDepartmentForHierarchy) -- Start with head of #1 dept
50
51     UNION ALL
52
53     SELECT
54         e_mgr.employeeId,
55         e_mgr.employeeName,
56         e_mgr.managerId,
57         eh.level + 1
58     FROM EmployeesIV e_mgr
59     JOIN EmployeeHierarchy eh ON e_mgr.employeeId = eh.managerId
60     WHERE eh.managerId IS NOT NULL -- Stop when managerId is NULL (top
61     of hierarchy)
62 )
63 SELECT
64     (SELECT 'Hierarchy for Head of Department: ' || tdH.departmentName
65     FROM TopDepartmentForHierarchy tdH) AS context,
66     eh.employeeId,
67     eh.employeeName,
68     eh.managerId,
69     eh.level
70 FROM EmployeeHierarchy eh
71 ORDER BY eh.level DESC;

```

Listing 15: Solution for Hardcore Problem (Part 2: Hierarchy for #1 Department Head)