

Ranking Functions in SQL: Solutions

Generated Examples

May 21, 2025

Contents

1	Practice Meanings, Values, Relations, and Advantages	3
1.1	Basic ROW_NUMBER()	3
1.2	Basic RANK()	3
1.3	Basic DENSE_RANK()	3
1.4	Comparing ROW_NUMBER(), RANK(), DENSE_RANK() within partitions	4
1.5	Advantage - Top N per group	4
2	Practice Disadvantages of Technical Concepts	5
2.1	Misinterpretation of RANK() vs DENSE_RANK() for Nth distinct value	5
2.2	Potential for confusion with complex ORDER BY in window definition .	6
2.3	Conceptual disadvantage - Readability with many window functions . . .	7
3	Practice Cases of Inefficient Alternatives	8
3.1	Find the employee(s) with the highest salary in each department	8
3.2	Assign sequential numbers to records	8
3.3	Ranking based on an aggregate	9
4	Practice Hardcore Combined Problem	11
4.1	Comprehensive Employee Analysis	11

Dataset for PostgreSQL

The following SQL code creates and populates the necessary tables for the exercises. (Same as in exercises document).

```
1 DROP TABLE IF EXISTS product\_sales;
2 DROP TABLE IF EXISTS employees;
3
4 CREATE TABLE employees (
5     employee\_id INT PRIMARY KEY,
6     first\_name VARCHAR(50),
7     last\_name VARCHAR(50),
8     department VARCHAR(50),
9     salary DECIMAL(10, 2),
10    hire\_date DATE
11 );
12
13 INSERT INTO employees (employee\_id, first\_name, last\_name, department, salary, hire\_
    _date) VALUES
14 (1, 'Alice', 'Smith', 'HR', 60000.00, '2020-01-15'),
15 (2, 'Bob', 'Johnson', 'HR', 65000.00, '2019-07-01'),
16 (3, 'Charlie', 'Williams', 'IT', 80000.00, '2021-03-10'),
17 (4, 'David', 'Brown', 'IT', 90000.00, '2018-05-20'),
18 (5, 'Eve', 'Jones', 'IT', 80000.00, '2022-01-05'),
19 (6, 'Frank', 'Garcia', 'Finance', 75000.00, '2020-11-01'),
20 (7, 'Grace', 'Miller', 'Finance', 75000.00, '2021-06-15'),
21 (8, 'Henry', 'Davis', 'Marketing', 70000.00, '2019-02-28'),
22 (9, 'Ivy', 'Rodriguez', 'Marketing', 72000.00, '2023-01-10'),
23 (10, 'Jack', 'Wilson', 'Marketing', 70000.00, '2020-08-15'),
24 (11, 'Karen', 'Moore', 'HR', 60000.00, '2021-09-01'),
25 (12, 'Liam', 'Taylor', 'IT', 95000.00, '2023-03-01');
26
27 CREATE TABLE product\_sales (
28     sale\_id INT PRIMARY KEY,
29     product\_name VARCHAR(100),
30     category VARCHAR(50),
31     sale\_date DATE,
32     sale\_amount DECIMAL(10, 2),
33     quantity\_sold INT
34 );
35
36 INSERT INTO product\_sales (sale\_id, product\_name, category, sale\_date, sale\_amount,
    quantity\_sold) VALUES
37 (1, 'Laptop Pro', 'Electronics', '2023-01-10', 1200.00, 5),
38 (2, 'Smartphone X', 'Electronics', '2023-01-12', 800.00, 10),
39 (3, 'Office Chair', 'Furniture', '2023-01-15', 150.00, 20),
40 (4, 'Desk Lamp', 'Furniture', '2023-01-18', 40.00, 30),
41 (5, 'Laptop Pro', 'Electronics', '2023-02-05', 1200.00, 3),
42 (6, 'Gaming Mouse', 'Electronics', '2023-02-10', 75.00, 50),
43 (7, 'Smartphone X', 'Electronics', '2023-02-15', 780.00, 8),
44 (8, 'Bookshelf', 'Furniture', '2023-02-20', 200.00, 10),
45 (9, 'Laptop Pro', 'Electronics', '2023-03-01', 1150.00, 4),
46 (10, 'External HDD', 'Electronics', '2023-03-05', 100.00, 25),
47 (11, 'Office Chair', 'Furniture', '2023-03-10', 140.00, 15),
48 (12, 'Desk Lamp', 'Furniture', '2023-03-15', 35.00, 40),
49 (13, 'Smartphone Y', 'Electronics', '2023-03-20', 900.00, 12),
50 (14, 'Coffee Maker', 'Appliances', '2023-01-20', 60.00, 10),
51 (15, 'Toaster', 'Appliances', '2023-02-25', 30.00, 15),
52 (16, 'Blender', 'Appliances', '2023-03-01', 50.00, 8);
```

Listing 1: Dataset for Employees and Product Sales Tables

1 Practice Meanings, Values, Relations, and Advantages

1.1 Basic ROW_NUMBER()

Problem: Assign a unique sequential number to each employee based on their hire date (oldest first).

Solution:

```
1 SELECT
2     employee\_id,
3     first\_name,
4     last\_name,
5     hire\_date,
6     ROW\_NUMBER() OVER (ORDER BY hire\_date ASC) as hiring\_sequence
7 FROM
8     employees;
```

Listing 2: Solution for i.1

1.2 Basic RANK()

Problem: Rank employees based on their salary in descending order. Show how ties are handled (gaps in rank).

Solution:

```
1 SELECT
2     employee\_id,
3     first\_name,
4     last\_name,
5     salary,
6     RANK() OVER (ORDER BY salary DESC) as salary\_rank
7 FROM
8     employees;
```

Listing 3: Solution for i.2

1.3 Basic DENSE_RANK()

Problem: Rank employees based on their salary in descending order using dense ranking. Show how ties are handled differently from RANK() (no gaps in rank).

Solution:

```
1 SELECT
2     employee\_id,
3     first\_name,
4     last\_name,
5     salary,
6     DENSE\_RANK() OVER (ORDER BY salary DESC) as dense\_salary\_rank
7 FROM
8     employees;
```

Listing 4: Solution for i.3

1.4 Comparing ROW_NUMBER(), RANK(), DENSE_RANK() within partitions

Problem: For each department, assign a unique row number, rank (with gaps), and dense rank (no gaps) to employees based on their salary in descending order.

Solution:

```
1 SELECT
2     employee\_id,
3     first\_name,
4     last\_name,
5     department,
6     salary,
7     ROW\_NUMBER() OVER (PARTITION BY department ORDER BY salary DESC)
8     as dept\_row\_num,
9     RANK() OVER (PARTITION BY department ORDER BY salary DESC) as dept\_
10    _salary\_rank,
11    DENSE\_RANK() OVER (PARTITION BY department ORDER BY salary DESC)
12    as dept\_dense\_salary\_rank
13 FROM
14     employees
15 ORDER BY
16     department, salary DESC;
```

Listing 5: Solution for i.4

1.5 Advantage - Top N per group

Problem: Identify the top 2 highest-paid employees in each department.

Solution:

```
1 WITH RankedEmployees AS (
2     SELECT
3         employee\_id,
4         first\_name,
5         last\_name,
6         department,
7         salary,
8         DENSE\_RANK() OVER (PARTITION BY department ORDER BY salary
9         DESC) as dept\_dense\_salary\_rank
10    FROM
11        employees
12 )
13 SELECT
14     employee\_id,
15     first\_name,
16     last\_name,
17     department,
18     salary,
19     dept\_dense\_salary\_rank
20 FROM
21     RankedEmployees
22 WHERE
23     dept\_dense\_salary\_rank <= 2
24 ORDER BY
25     department, dept\_dense\_salary\_rank;
```

Listing 6: Solution for i.5

2 Practice Disadvantages of Technical Concepts

2.1 Misinterpretation of RANK() vs DENSE_RANK() for Nth distinct value

Problem: A manager wants to identify all employees who are in the top 2 distinct salary tiers within the company. Demonstrate this by comparing results from RANK() and DENSE_RANK(). Specifically, if the requirement is "employees in the 4th highest salary tier company-wide", show how using RANK() = 4 might yield no results, while DENSE_RANK() = 4 would correctly identify them.

Solution: First, let's display company-wide ranks using both functions:

```
1 SELECT
2     employee\_id, first\_name, salary,
3     RANK() OVER (ORDER BY salary DESC) as company\_rank,
4     DENSE\_RANK() OVER (ORDER BY salary DESC) as company\_dense\_rank
5 FROM employees
6 ORDER BY salary DESC;
```

Listing 7: Solution for ii.1 - Display Ranks

Observation from the above query output:

- Distinct salaries are: 95k, 90k, 80k, 75k, 72k, 70k, 65k, 60k.
- DENSE_RANK correctly numbers these distinct salary tiers from 1 to 8.
- RANK results:
 - Liam (95k) -> company_rank=1, company_dense_rank=1
 - David (90k) -> company_rank=2, company_dense_rank=2
 - Charlie (80k), Eve (80k) -> Both company_rank=3, company_dense_rank=3
 - Frank (75k), Grace (75k) -> Both company_rank=5, company_dense_rank=4
 - Ivy (72k) -> company_rank=7, company_dense_rank=5
 - Henry (70k), Jack (70k) -> Both company_rank=8, company_dense_rank=6
 - Bob (65k) -> company_rank=10, company_dense_rank=7
 - Alice (60k), Karen (60k) -> Both company_rank=11, company_dense_rank=8

If a manager asks for "employees in the 4th highest salary tier", using company_rank = 4 would yield no results from the above. Using company_dense_rank = 4 correctly identifies Frank and Grace (salary 75000.00).

Query to find employees in the 4th distinct salary tier using DENSE_RANK():

```
1 SELECT
2     employee\_id, first\_name, salary,
3     DENSE\_RANK() OVER (ORDER BY salary DESC) as company\_dense\_rank
4 FROM employees
5 QUALIFY DENSE\_RANK() OVER (ORDER BY salary DESC) = 4;
6 -- Using subquery for broader compatibility:
7 -- WITH RankedSalaries AS (
8 --     SELECT employee\_id, first\_name, salary,
9 --     DENSE\_RANK() OVER (ORDER BY salary DESC) as company\_dense\_rank
10 -- FROM employees
```

```

11 -- )
12 -- SELECT * FROM RankedSalaries WHERE company\_dense\_rank = 4;

```

Listing 8: Solution for ii.1 - Correct identification

This demonstrates a disadvantage if `RANK()` is chosen when `DENSE_RANK()` is appropriate for the business question regarding distinct value tiers.

2.2 Potential for confusion with complex ORDER BY in window definition

Problem: Employees are ranked within their department by salary (descending). For employees with the same salary, their secondary sort key is hire date (ascending, earlier hire gets better rank). Show this ranking using `ROW_NUMBER()`. Then, demonstrate how if the secondary sort key (`hire_date`) was mistakenly ordered descending, it would change the row numbers for tied-salary employees, potentially leading to incorrect "top" employee identification if the secondary sort was critical for tie-breaking. Focus on employees in departments and salary groups where ties exist.

Solution:

```

1 -- Correct secondary sort (hire\_date ASC for tie-breaking)
2 SELECT
3     department, first\_name, salary, hire\_date,
4     ROW\_NUMBER() OVER (PARTITION BY department ORDER BY salary DESC,
5                          hire\_date ASC) as rn\_hire\_asc
6 FROM employees
7 WHERE salary IN (80000, 75000, 70000, 60000) -- Focus on employees with
8                                             tied salaries
9 ORDER BY department, salary DESC, hire\_date ASC;
10
11 -- Incorrect/different secondary sort (hire\_date DESC for tie-breaking)
12 SELECT
13     department, first\_name, salary, hire\_date,
14     ROW\_NUMBER() OVER (PARTITION BY department ORDER BY salary DESC,
15                          hire\_date DESC) as rn\_hire\_desc
16 FROM employees
17 WHERE salary IN (80000, 75000, 70000, 60000) -- Focus on employees with
18                                             tied salaries
19 ORDER BY department, salary DESC, hire\_date DESC;

```

Listing 9: Solution for ii.2 - Comparing `ROW_NUMBER()` with different secondary sort

Observation:

- In IT, for salary 80000.00: Charlie (hired 2021-03-10), Eve (hired 2022-01-05).
 - With `rn_hire_asc`: Charlie is 1st among 80k, Eve is 2nd. (After David 90k and Liam 95k in IT, so actual ranks are 3rd and 4th if considering top salary). The problem stated "salary (descending)", so within IT: Liam (95k) is 1, David (90k) is 2. Then for 80k: Charlie (hire 2021) is 3, Eve (hire 2022) is 4.
 - With `rn_hire_desc`: For 80k in IT: Eve (hire 2022, later) becomes 3, Charlie (hire 2021, earlier) becomes 4.
- Similar changes occur in HR for 60000.00 (Alice vs Karen) and Finance for 75000.00 (Frank vs Grace), and Marketing for 70000.00 (Henry vs Jack).

This shows how a subtle change in the `ORDER BY` clause within the window definition can significantly alter rankings for records that are tied on the primary sort key. This can be a source of confusion or error if the tie-breaking rule is critical and not correctly implemented or understood.

2.3 Conceptual disadvantage - Readability with many window functions

Problem: Display each employee's salary, their salary rank within their department, their overall salary rank in the company (dense), and a row number based on alphabetical order of their full name (`last_name`, then `first_name`). The query itself will demonstrate how multiple window functions can make the `SELECT` clause dense.

Solution:

```
1 SELECT
2     e.employee_id,
3     e.first_name,
4     e.last_name,
5     e.department,
6     e.salary,
7     RANK() OVER (PARTITION BY e.department ORDER BY e.salary DESC) as
      dept_salary_rank,
8     DENSE_RANK() OVER (ORDER BY e.salary DESC) as overall_company_
      _dense_salary_rank,
9     ROW_NUMBER() OVER (ORDER BY e.last_name ASC, e.first_name ASC)
      as alphabetical_row_num
10 FROM
11     employees e
12 ORDER BY
13     e.department, dept_salary_rank;
```

Listing 10: Solution for ii.3

Observation: The `SELECT` clause, while powerful, becomes somewhat lengthy with multiple window functions, each with its own `OVER()` clause. In much more complex queries involving joins, aggregations, and even more window functions with varying partitions and orderings, this part of the query can become very difficult to read, debug, and maintain. This highlights a "soft" disadvantage related to code complexity and readability. Using CTEs can help mitigate this but doesn't entirely remove the conceptual load of tracking multiple window contexts.

3 Practice Cases of Inefficient Alternatives

3.1 Find the employee(s) with the highest salary in each department

Problem: List the employee(s) with the highest salary in each department. If there are ties, list all. Provide both an inefficient solution (e.g., using a subquery with `MAX` for each department and joining back) and an efficient solution using ranking functions.

Solution: *Inefficient Solution (using subquery with `MAX` for each department):*

```
1 SELECT e.department, e.first\_name, e.last\_name, e.salary
2 FROM employees e
3 JOIN (
4     SELECT department, MAX(salary) AS max\_salary
5     FROM employees
6     GROUP BY department
7 ) AS dept\_max\_salary
8 ON e.department = dept\_max\_salary.department AND e.salary = dept\_max\_salary.max\_salary
9 ORDER BY e.department, e.salary DESC;
```

Listing 11: Inefficient Solution for iii.1

Efficient Solution (using `DENSE_RANK()`):

```
1 WITH RankedSalaries AS (
2     SELECT
3         department,
4         first\_name,
5         last\_name,
6         salary,
7         DENSE\_RANK() OVER (PARTITION BY department ORDER BY salary
8         DESC) as rn
9     FROM employees
10 )
11 SELECT department, first\_name, last\_name, salary
12 FROM RankedSalaries
13 WHERE rn = 1
14 ORDER BY department, salary DESC;
```

Listing 12: Efficient Solution for iii.1

Note: The efficient solution scans the table once to compute ranks, while the inefficient one typically involves an aggregation and then a join, which can be less performant on large datasets, especially without proper indexing.

3.2 Assign sequential numbers to records

Problem: Provide a unique sequential number for each product sale, ordered by `sale_date` then by `sale_id`. Provide both an inefficient solution (e.g., using a correlated subquery to count preceding records) and an efficient solution using ranking functions.

Solution: *Inefficient Solution (using a correlated subquery):*

```
1 SELECT
2     s1.sale\_id,
3     s1.product\_name,
4     s1.sale\_date,
5     (SELECT COUNT(*)
```



```

6      FROM product\_sales s2
7      WHERE (s2.sale\_date < s1.sale\_date)
8             OR (s2.sale\_date = s1.sale\_date AND s2.sale\_id <= s1.sale\_
9             _id)
10     ) as inefficient\_row\_num
11 FROM product\_sales s1
12 ORDER BY inefficient\_row\_num;

```

Listing 13: Inefficient Solution for iii.2

Efficient Solution (using ROW_NUMBER()):

```

1 SELECT
2     sale\_id,
3     product\_name,
4     sale\_date,
5     ROW\_NUMBER() OVER (ORDER BY sale\_date ASC, sale\_id ASC) as
6     efficient\_row\_num
7 FROM product\_sales
8 ORDER BY efficient\_row\_num;

```

Listing 14: Efficient Solution for iii.2

Note: Correlated subqueries like the one in the inefficient solution are often very slow because the subquery is executed for each row of the outer query. ROW_NUMBER() is optimized for this task.

3.3 Ranking based on an aggregate

Problem: Rank product categories by their total sales amount in descending order. Provide both an inefficient solution (e.g., aggregating in a subquery/CTE, then using another correlated subquery or complex logic for ranking) and an efficient solution using ranking functions on the aggregated results.

Solution: *Inefficient Solution (aggregating then using a correlated subquery for ranking):*

```

1 WITH CategorySales AS (
2     SELECT
3         category,
4         SUM(sale\_amount) AS total\_sales
5     FROM product\_sales
6     GROUP BY category
7 )
8 SELECT
9     cs.category,
10    cs.total\_sales,
11    (SELECT COUNT(DISTINCT cs2.total\_sales)
12     FROM CategorySales cs2
13     WHERE cs2.total\_sales >= cs.total\_sales) AS inefficient\_dense\_
14    _rank
15 FROM CategorySales cs
16 ORDER BY inefficient\_dense\_rank;

```

Listing 15: Inefficient Solution for iii.3

Efficient Solution (aggregating then using a ranking function):

```

1 WITH CategorySales AS (
2     SELECT

```

```

3         category,
4         SUM(sale\_amount) AS total\_sales
5     FROM product\_sales
6     GROUP BY category
7 )
8 SELECT
9     category,
10    total\_sales,
11    DENSE\_RANK() OVER (ORDER BY total\_sales DESC) as efficient\_sales
12    \_rank
13 FROM CategorySales
14 ORDER BY efficient\_sales\_rank;

```

Listing 16: Efficient Solution for iii.3

Note: Similar to the previous example, the correlated subquery for ranking in the inefficient solution performs poorly compared to the direct application of a window ranking function on the aggregated results.

4 Practice Hardcore Combined Problem

4.1 Comprehensive Employee Analysis

Problem: For each department:

1. Identify the employee(s) with the highest salary in that department.
2. For these top-paid employees, determine their salary rank across the entire company (use RANK() for ties).
3. Show the average salary of their department.
4. Calculate the difference between their salary and their department's average salary.
5. Identify the `employee_id` and full name (`first_name || ' ' || last_name`) of the employee who was hired immediately *after* them within the same department, along with that next hire's date. If they are the last one hired in their department (among the top-paid, or overall if simpler for this part), this should be NULL for the next hire's details. (Clarification: "last one hired in their department" should be interpreted as the one with latest hire date in the entire department for identifying the "next hire").

Present a single SQL query that provides all this information.

Solution:

```
1 WITH DepartmentTopSalary AS (  
2     SELECT  
3         employee_id,  
4         first_name,  
5         last_name,  
6         department,  
7         salary,  
8         hire_date,  
9         DENSE_RANK() OVER (PARTITION BY department ORDER BY salary DESC) as dept_salary_rank_val, -- Renamed to  
10        avoid_conflict  
11        AVG(salary) OVER (PARTITION BY department) as avg_dept_salary  
12    FROM employees  
13 ),  
14 TopPaidInDepartment AS (  
15     SELECT  
16         employee_id,  
17         first_name,  
18         last_name,  
19         department,  
20         salary,  
21         hire_date,  
22         avg_dept_salary,  
23         RANK() OVER (ORDER BY salary DESC) as overall_company_salary_rank  
24    FROM DepartmentTopSalary  
25   WHERE dept_salary_rank_val = 1 -- Used renamed column  
26 ),  
27 EmployeeHireSequence AS (  
28     SELECT  
29         employee_id,  
30         first_name, -- Added for next_hire full name  
31         last_name,  -- Added for next_hire full name  
32         department,  
33         hire_date,  
34         -- This ROW_NUMBER is for all employees in the department by hire_date  
35         -- to correctly identify the "next hire" overall in the department.  
36         ROW_NUMBER() OVER (PARTITION BY department ORDER BY hire_date ASC) as hire_sequence_in_dept  
37    FROM employees  
38 ),  
39 SELECT  
40     tpid.department,  
41     tpid.first_name AS top_emp_first_name,  
42     tpid.last_name AS top_emp_last_name,  
43     tpid.salary AS top_emp_salary,  
44     tpid.overall_company_salary_rank,  
45     ROUND(tpid.avg_dept_salary, 2) AS department_avg_salary,  
46     ROUND(tpid.salary - tpid.avg_dept_salary, 2) AS salary_diff_from_dept_avg,  
47     tpid.hire_date AS top_emp_hire_date,  
48     next_hire.employee_id AS next_hired_employee_id,  
49     next_hire.first_name || ' ' || next_hire.last_name AS next_hired_full_name,  
50     next_hire.hire_date AS next_hired_hire_date  
51 FROM TopPaidInDepartment tpid  
52 -- Join with EmployeeHireSequence to get the hire_sequence_in_dept for the top paid employee  
53 JOIN EmployeeHireSequence current_emp_hire_seq
```

```

53 ON tpid.employee\_id = current\_emp\_hire\_seq.employee\_id
54 -- LEFT JOIN to find the next employee in the sequence within the same department
55 LEFT JOIN EmployeeHireSequence next\_hire
56 ON tpid.department = next\_hire.department -- Ensure same department
57 AND current\_emp\_hire\_seq.hire\_sequence\_in\_dept + 1 = next\_hire.hire\_sequence\_in\_dept
58 ORDER BY
59 tpid.department, tpid.salary DESC, tpid.employee\_id;

```

Listing 17: Solution for iv.1

Explanation of the Hardcore Problem Solution Logic:

- DepartmentTopSalary CTE:
 - Uses DENSE_RANK() partitioned by department and ordered by salary DESC to find the rank of each employee's salary within their department.
 - Uses AVG(salary) as a window aggregate function partitioned by department to calculate the average salary for each department and associate it with every employee in that department.
- TopPaidInDepartment CTE:
 - Filters DepartmentTopSalary for employees where dept_salary_rank_val = 1, effectively selecting the highest-paid employee(s) in each department.
 - Calculates overall_company_salary_rank for these top-paid employees using RANK() ordered by salary DESC across the entire company.
- EmployeeHireSequence CTE:
 - Assigns a ROW_NUMBER() to each employee within their department based on their hire_date (earliest first). This sequence is crucial for identifying the "next hired" employee.
- Final SELECT Statement:
 - Starts with the TopPaidInDepartment employees.
 - Joins with EmployeeHireSequence (aliased as current_emp_hire_seq) on employee_id to get the hire_sequence_in_dept of the current top-paid employee.
 - LEFT JOINS again with EmployeeHireSequence (aliased as next_hire) on two conditions:
 1. Same department.
 2. The hire_sequence_in_dept of next_hire is exactly one greater than that of current_emp_hire_seq. This finds the employee hired immediately after the current top-paid employee in that department.

A LEFT JOIN is used because if a top-paid employee is the last one hired in their department, there will be no "next hire," and their details should be NULL.
 - Selects the required fields, including calculated difference and formatting the average salary.
 - Orders the results for clarity.