

Query Optimization and Performance

Indexing Strategies, EXPLAIN Plans, Optimizing Window Functions and Aggregates: *Solutions*

Sequential SQL

May 21, 2025

Contents

1	Solutions for Indexing Strategies	2
1.1	Solution for Exercise IS-1	2
1.2	Solution for Exercise IS-2	2
1.3	Solution for Exercise IS-3	3
1.4	Solution for Exercise IS-4	3
1.5	Solution for Exercise IS-5	5
1.6	Solution for Exercise IS-6	6
1.7	Solution for Exercise IS-7	7
2	Solutions for EXPLAIN Plans	8
2.1	Solution for Exercise EP-1	8
2.2	Solution for Exercise EP-2	8
2.3	Solution for Exercise EP-3	9
2.4	Solution for Exercise EP-4	10
3	Solutions for Optimizing Window Functions and Aggregates	10
3.1	Solution for Exercise OWA-1	10
3.2	Solution for Exercise OWA-2	11
3.3	Solution for Exercise OWA-3	11
3.4	Solution for Exercise OWA-4	12

1 Solutions for Indexing Strategies

1.1 Solution for Exercise IS-1

1. Query:

```
1 SELECT employeeId, firstName, lastName, jobTitle FROM Employees
   WHERE jobTitle = 'Data Analyst';
```

2. Observe performance (before specific index on jobTitle):

```
1 EXPLAIN ANALYZE SELECT employeeId, firstName, lastName, jobTitle
  FROM Employees WHERE jobTitle = 'Data Analyst';
2 -- Expected: Seq Scan on Employees if no specific index on jobTitle
   exists or is chosen.
```

3. Create B-tree index:

```
1 CREATE INDEX idxEmployeesJobTitle ON Employees (jobTitle);
```

4. Re-run and observe improvement:

```
1 EXPLAIN ANALYZE SELECT employeeId, firstName, lastName, jobTitle
  FROM Employees WHERE jobTitle = 'Data Analyst';
2 -- Expected: Index Scan or Bitmap Heap Scan using
   idxEmployeesJobTitle.
```

Description: The plan changes from a Sequential Scan (reading all rows) to an Index-based scan. The B-tree index on `jobTitle` allows the database to quickly locate an ordered list of pointers to rows where `jobTitle` is 'Data Analyst', significantly reducing I/O and CPU by avoiding a full table scan. This is advantageous for equality predicates on moderately to highly selective columns.

1.2 Solution for Exercise IS-2

1. **Write Overhead Description:** With 10 indexes, each INSERT operation would require updating all 10 index structures in addition to writing the row to the table itself. With only 2 indexes, only those 2 index structures need updating. Each index update involves finding the correct position in the B-tree (or other structure) and inserting the new key/pointer, potentially causing page splits and other maintenance. Therefore, INSERT performance would be noticeably slower with 10 indexes compared to 2 indexes. **Disadvantage Illustrated:** Indexes add overhead to DML operations (INSERT, UPDATE, DELETE). Over-indexing can significantly degrade write performance.

2. **Very Low Selectivity / Small Table:**

```
1 CREATE INDEX IF NOT EXISTS idxDepartmentsLocation ON Departments (
  location);
2 EXPLAIN ANALYZE SELECT departmentName FROM Departments WHERE
  location = 'Building A, Floor 1';
```

Observation & Explanation: The optimizer might still choose a Seq Scan. For very small tables, the cost of reading the entire table (which might be just one or a few data blocks) can be less than the overhead of an index scan (reading index

blocks then data blocks via pointers). If 'Building A, Floor 1' is very common (low selectivity), fetching many rows via an index (random I/Os) can be costlier than one sequential scan. The optimizer compares estimated costs.

3. Leading Wildcard LIKE:

```
1 EXPLAIN ANALYZE SELECT employeeId, email FROM Employees WHERE email
   LIKE '%user123%';
```

Observation & Explanation: The B-tree index on `email` will likely **not** be used effectively (or at all) for the `WHERE email LIKE '%user123%'` condition, resulting in a `Seq Scan`. B-tree indexes are structured based on the prefix of the indexed values. A leading wildcard (%) means the beginning of the string is unknown, so the database cannot use the B-tree's ordered structure to efficiently narrow down the search. It must check all values. (Note: For `email LIKE 'user123%'`, the index would be effective).

1.3 Solution for Exercise IS-3

1. Naive LIKE query:

```
1 EXPLAIN ANALYZE SELECT projectId, projectName FROM Projects
2 WHERE projectDescription LIKE '%innovation%' AND projectDescription
   LIKE '%strategy%';
3 -- Expected: Seq Scan on Projects, slow due to full text scanning
   for each LIKE.
```

2. Create GIN index:

```
1 DROP INDEX IF EXISTS idxGinProjectDescription;
2 CREATE INDEX idxGinProjectDescription ON Projects USING GIN (
   to_tsvector('english', projectDescription));
```

3. FTS query:

```
1 EXPLAIN ANALYZE SELECT projectId, projectName FROM Projects
2 WHERE to_tsvector('english', projectDescription) @@ to_tsquery('
   english', 'innovation & strategy');
```

4. **Comparison and GIN Advantage:** The FTS query using the GIN index will be significantly faster and use a plan like `Bitmap Heap Scan on Projects` with a `Bitmap Index Scan on idxGinProjectDescription`. GIN (Generalized Inverted Index) is designed for indexing composite values where elements within the value are of interest (like words in text). `to_tsvector` breaks text into lexemes. The GIN index stores these lexemes and lists of documents they appear in. **Advantage:** Instead of scanning all text, FTS quickly finds documents containing "innovation" and documents containing "strategy" using the index, then finds the intersection. This is vastly more efficient than B-trees (unsuitable for '%text%') or repeated full string scans with LIKE.

1.4 Solution for Exercise IS-4

1. SQL Query:

```

1 WITH EligibleEmployees AS (
2     SELECT
3         e.employeeId,
4         e.firstName,
5         e.lastName,
6         e.jobTitle,
7         e.hireDate,
8         e.departmentId
9     FROM
10        Employees e
11    WHERE
12        e.status = 'Active'
13        AND e.departmentId IN (SELECT departmentId FROM Departments
14                                WHERE departmentName IN ('Engineering', 'Product Management'))
15        AND e.hireDate BETWEEN '2015-01-01' AND '2020-12-31'
16        AND e.performanceScore >= 3.5
17 ),
18 EmployeeProjectCounts AS (
19     SELECT
20         ee.employeeId,
21         ee.firstName,
22         ee.lastName,
23         ee.jobTitle,
24         ee.hireDate,
25         d.departmentName,
26         COUNT(ep.projectId) AS projectCount
27     FROM
28        EligibleEmployees ee
29     JOIN
30        Departments d ON ee.departmentId = d.departmentId
31     LEFT JOIN
32        EmployeeProjects ep ON ee.employeeId = ep.employeeId
33     GROUP BY
34         ee.employeeId, ee.firstName, ee.lastName, ee.jobTitle, ee.
35         hireDate, d.departmentName
36 )
37 SELECT
38     firstName,
39     lastName,
40     jobTitle,
41     departmentName,
42     hireDate,
43     projectCount
44 FROM
45     EmployeeProjectCounts
46 ORDER BY
47     departmentName ASC, projectCount DESC, hireDate DESC;

```

2. Candidate Columns for Indexing:

- Employees table: status (WHERE), departmentId (WHERE, JOIN), hireDate (WHERE, ORDER BY), performanceScore (WHERE), employeeId (JOIN, GROUP BY - PK).
- Departments table: departmentId (JOIN - PK), departmentName (WHERE subquery, ORDER BY - UNIQUE).
- EmployeeProjects table: employeeId (JOIN, COUNT), projectId (COUNT).

3. Proposed B-tree Indexes & Choices:

- On Employees:
 - CREATE INDEX IF NOT EXISTS idxEmployeesStatus ON Employees (status);
 - CREATE INDEX IF NOT EXISTS idxEmployeesPerformanceScore ON Employees (performanceScore);
 - (Ensure idxEmployeesHireDate ON Employees(hireDate) exists)
 - (Ensure idxEmployeesDepartmentId ON Employees(departmentId) exists)
- On Departments: (PK departmentId and UNIQUE departmentName are already indexed.)
- On EmployeeProjects:
 - CREATE INDEX IF NOT EXISTS idxEmployeeProjectsEmpId ON EmployeeProjects (employeeId);

4. Create indexes and Analyze:

```
1 CREATE INDEX IF NOT EXISTS idxEmployeesStatus ON Employees (status)
;
2 CREATE INDEX IF NOT EXISTS idxEmployeesPerformanceScore ON
  Employees (performanceScore);
3 CREATE INDEX IF NOT EXISTS idxEmployeeProjectsEmpId ON
  EmployeeProjects (employeeId);
4
5 EXPLAIN ANALYZE
6 WITH EligibleEmployees AS ( /* ...as above... */ ),
7 EmployeeProjectCounts AS ( /* ...as above... */ )
8 SELECT firstName, lastName, jobTitle, departmentName, hireDate,
   projectCount
9 FROM EmployeeProjectCounts
10 ORDER BY departmentName ASC, projectCount DESC, hireDate DESC;
```

Conceptual Plan Description:

- Subquery for department IDs: Fast Index Scan on Departments.departmentName.
- EligibleEmployees CTE: Likely a BitmapAnd/BitmapOr operation combining results from multiple index scans on Employees (on status, departmentId, hireDate, performanceScore), followed by a Bitmap Heap Scan.
- EmployeeProjectCounts CTE: Join to Departments, Left Join to EmployeeProjects (using idxEmployeeProjectsEmpId), then HashAggregate.
- Final SELECT: A Sort operation.

The indexes significantly reduce rows scanned from Employees, making filtering efficient. The index on EmployeeProjects.employeeId speeds up counting projects.

1.5 Solution for Exercise IS-5

1. Query:

```
1 SELECT projectId, projectName, startDate FROM Projects WHERE
   startDate > '2023-01-01';
```

2. Observe performance (before BRIN index):

```
1 EXPLAIN ANALYZE SELECT projectId, projectName, startDate FROM
  Projects WHERE startDate > '2023-01-01';
2 -- Expected: Seq Scan on Projects, as no specific index on
  startDate exists in the initial setup.
```

3. Create BRIN index:

```
1 CREATE INDEX idxProjectsStartDateBrin ON Projects USING BRIN (
  startDate);
```

4. Re-run and observe improvement:

```
1 EXPLAIN ANALYZE SELECT projectId, projectName, startDate FROM
  Projects WHERE startDate > '2023-01-01';
2 -- Expected: Bitmap Heap Scan on Projects with Bitmap Index Scan on
  idxProjectsStartDateBrin.
```

Description: The plan shifts from a Seq Scan to a Bitmap Heap Scan using the BRIN index. BRIN (Block Range Index) is ideal for large tables with naturally ordered columns like `startDate`, which tends to be sequential due to project creation over time. BRIN stores summaries of data ranges (min/max values per block), allowing the database to skip irrelevant blocks. This reduces I/O significantly compared to scanning the entire table, especially for range queries on large datasets. BRIN's small size and low maintenance overhead make it advantageous over B-tree indexes for such scenarios, though it's less precise and best for sequential data.

1.6 Solution for Exercise IS-6

1. Query:

```
1 SELECT employeeId, firstName, lastName, email FROM Employees WHERE
  email = 'user100@example.com';
```

2. Observe performance (with B-tree index):

```
1 EXPLAIN ANALYZE SELECT employeeId, firstName, lastName, email FROM
  Employees WHERE email = 'user100@example.com';
2 -- Expected: Index Scan using the UNIQUE B-tree index on email (
  from UNIQUE constraint).
```

3. Drop UNIQUE constraint, create Hash index, and restore constraint:

```
1 -- Drop the UNIQUE constraint (which drops its B-tree index)
2 ALTER TABLE Employees DROP CONSTRAINT employees_email_key;
3
4 -- Create Hash index CONCURRENTLY
5 CREATE INDEX CONCURRENTLY idxEmployeesEmailHash ON Employees USING
  HASH (email);
6
7 -- Restore UNIQUE constraint (creates a new B-tree index for
  constraint enforcement)
8 ALTER TABLE Employees ADD CONSTRAINT employees_email_key UNIQUE (
  email);
```

4. Re-run and analyze:

```
1 EXPLAIN ANALYZE SELECT employeeId, firstName, lastName, email FROM
  Employees WHERE email = 'user100@example.com';
2 -- Expected: Likely Index Scan using the UNIQUE B-tree index (from
  restored constraint), not the Hash index.
```

Explanation: The query planner may prefer the B-tree index created by the UNIQUE constraint over the Hash index, as B-trees are versatile and reliable for equality lookups, and the constraint ensures uniqueness. However, Hash indexes are optimized specifically for equality comparisons (=), storing hashed values for faster lookups with a smaller footprint than B-trees. They lack support for range queries or sorting, but for exact matches, they can be slightly faster due to simpler structure. Using CONCURRENTLY is critical in production, as it allows index creation without locking the table, preventing downtime during writes. The planner's choice depends on cost estimates, but Hash indexes shine in high-read, equality-heavy workloads.

1.7 Solution for Exercise IS-7

1. Add generated tsvector column:

```
1 ALTER TABLE Projects
2 ADD COLUMN descriptionTsv TSVECTOR
3 GENERATED ALWAYS AS (to_tsvector('english', projectDescription))
  STORED;
```

2. Create GIN index with INCLUDE:

```
1 CREATE INDEX idxProjectsDescriptionTsvGin ON Projects USING GIN (
  descriptionTsv) INCLUDE (projectName, startDate);
```

3. Query for full-text search:

```
1 SELECT projectName, startDate FROM Projects
2 WHERE descriptionTsv @@ to_tsquery('english', 'agile & release');
```

4. Analyze performance:

```
1 EXPLAIN ANALYZE SELECT projectName, startDate FROM Projects
2 WHERE descriptionTsv @@ to_tsquery('english', 'agile & release');
3 -- Expected: Index Only Scan using idxProjectsDescriptionTsvGin.
```

Explanation: The plan uses an Index Only Scan, retrieving projectName and startDate directly from the GIN index without accessing the table heap. The stored tsvector column precomputes the lexemes from projectDescription, avoiding runtime computation of to_tsvector, which speeds up query execution. The GIN index efficiently handles full-text search by mapping lexemes to rows, and the INCLUDE clause adds projectName and startDate to the index, enabling a covering index. This eliminates heap access, reducing I/O and boosting performance compared to a regular GIN index on to_tsvector('english', projectDescription), which would require table access for non-indexed columns.

2 Solutions for EXPLAIN Plans

2.1 Solution for Exercise EP-1

1. Query:

```
1 SELECT
2     e.firstName, e.lastName, e.jobTitle, d.departmentName
3 FROM
4     Employees e
5 JOIN
6     Departments d ON e.departmentId = d.departmentId
7 WHERE
8     d.departmentName = 'Sales';
```

2. EXPLAIN output analysis:

```
1 EXPLAIN
2 SELECT
3     e.firstName, e.lastName, e.jobTitle, d.departmentName
4 FROM
5     Employees e
6 JOIN
7     Departments d ON e.departmentId = d.departmentId
8 WHERE
9     d.departmentName = 'Sales';
```

Identify:

- Scan type on Employees: Likely Index Scan on idxEmployeesDepartmentId.
- Scan type on Departments: Index Scan on departmentName (UNIQUE index).
- Join type: Likely Hash Join or Nested Loop.

3. "cost", "rows", "width" explanation:

- cost=X..Y: Estimated startup (X) and total (Y) cost in arbitrary units.
- rows=N: Estimated number of rows output by the node.
- width=W: Estimated average row width in bytes.

2.2 Solution for Exercise EP-2

1. Query: SELECT * FROM Employees WHERE salary > 150000;

2. EXPLAIN before insert:

```
1 EXPLAIN SELECT * FROM Employees WHERE salary > 150000;
```

3. Insert a high earner:

```
1 INSERT INTO Employees (firstName, lastName, email, departmentId,
2     salary, hireDate, jobTitle, performanceScore, status)
3 VALUES ('High', 'Earner', 'high.earner@example.com',
4     (SELECT departmentId FROM Departments WHERE departmentName
5     = 'Finance' LIMIT 1),
6     200000.00, CURRENT_DATE, 'CFO', 5.0, 'Active');
```


4. Re-run EXPLAIN:

```
1 EXPLAIN SELECT * FROM Employees WHERE salary > 150000;
```

Estimated rows won't change significantly without ANALYZE, as statistics are stale. This shows EXPLAIN's reliance on outdated stats.

5. Run EXPLAIN ANALYZE:

```
1 EXPLAIN ANALYZE SELECT * FROM Employees WHERE salary > 150000;
```

Comparison: actual time and actual rows reflect real execution, revealing discrepancies due to stale stats. ANALYZE provides actual performance data.

2.3 Solution for Exercise EP-3

1. Correlated Subquery:

```
1 EXPLAIN ANALYZE
2 SELECT
3     e.employeeId,
4     e.firstName,
5     e.lastName,
6     (SELECT p.projectName
7      FROM Projects p
8      JOIN EmployeeProjects epFind ON p.projectId = epFind.projectId
9      WHERE epFind.employeeId = e.employeeId AND p.projectName = '
10    Project Alpha 1'
11    LIMIT 1) AS projectAlphaName
12 FROM
13     Employees e
14 WHERE EXISTS (
15     SELECT 1 FROM EmployeeProjects epChk
16     JOIN Projects pChk ON epChk.projectId = pChk.projectId
17     WHERE epChk.employeeId = e.employeeId AND pChk.projectName = '
18     Project Alpha 1'
19 );
```

2. Analyze: The subquery runs per row, causing high loop counts and slow performance.

3. Rewrite using JOIN:

```
1 EXPLAIN ANALYZE
2 SELECT
3     e.employeeId, e.firstName, e.lastName, p.projectName
4 FROM
5     Employees e
6 JOIN
7     EmployeeProjects ep ON e.employeeId = ep.employeeId
8 JOIN
9     Projects p ON ep.projectId = p.projectId
10 WHERE
11     p.projectName = 'Project Alpha 1';
```

4. Comparison: The JOIN version uses efficient set-based operations (e.g., Hash Join), scanning tables once, making it faster than row-by-row subquery execution.

2.4 Solution for Exercise EP-4

```
1 SELECT
2     d.departmentName ,
3     COUNT(e.employeeId) as numEngineers ,
4     AVG(e.salary) as avgSalary
5 FROM
6     Departments d
7 JOIN
8     Employees e ON d.departmentId = e.departmentId
9 WHERE
10    e.jobTitle = 'Software Engineer'
11    AND e.hireDate > '2018-01-01'
12 GROUP BY
13     d.departmentId, d.departmentName
14 HAVING
15     AVG(e.salary) > 75000
16 ORDER BY
17     avgSalary DESC;
```

1. Run EXPLAIN (ANALYZE, BUFFERS):

```
1 EXPLAIN (ANALYZE , BUFFERS) SELECT /* ... as above ... */;
```

2. Time-consuming operations: Likely Seq Scan on Employees or Sort for ORDER BY.
3. Discrepancies: Large differences between estimated and actual rows suggest stale statistics.
4. High shared read: Indicates heavy disk I/O, likely from Seq Scan.
5. Improvements:
 - Composite index: CREATE INDEX idxEmpJobHireDept ON Employees (jobTitle, hireDate, departmentId) INCLUDE (salary);
 - Increase work_mem: SET LOCAL work_mem = '128MB'; to reduce disk spills.

3 Solutions for Optimizing Window Functions and Aggregates

3.1 Solution for Exercise OWA-1

1. Query:

```
1 SELECT
2     st.transactionId ,
3     st.customerId ,
4     st.transactionDate ,
5     st.totalAmount ,
6     AVG(st.totalAmount) OVER (PARTITION BY st.customerId) AS
7     avgCustomerSalesAmount
8 FROM
9     SalesTransactions st
10 ORDER BY
11     st.customerId, st.transactionDate
12 LIMIT 100;
```

2. Advantage: Window functions are concise, optimized for single-pass processing, and preserve row details unlike GROUP BY.

3.2 Solution for Exercise OWA-2

1. Query:

```
1 EXPLAIN ANALYZE
2 SELECT
3     transactionId,
4     totalAmount,
5     RANK() OVER (ORDER BY totalAmount DESC) AS overallRank
6 FROM
7     SalesTransactions
8 LIMIT 100;
```

2. Disadvantage: Sorting 1.5M rows is resource-intensive, risking disk spills.
3. With PARTITION BY productId: Smaller partitions reduce sort overhead and enable parallelism.

3.3 Solution for Exercise OWA-3

1. Inefficient Sketch: Correlated subquery per customer-month is slow due to repeated scans.

2. Optimized Query:

```
1 WITH CustomerMonthlySales2022 AS (
2     SELECT
3         st.customerId,
4         DATE_TRUNC('month', st.transactionDate) AS saleMonth,
5         SUM(st.totalAmount) AS monthlySales
6     FROM
7         SalesTransactions st
8     WHERE
9         st.transactionDate >= '2022-01-01' AND st.transactionDate <
10        '2023-01-01'
11    GROUP BY
12        st.customerId, DATE_TRUNC('month', st.transactionDate)
13 )
14 SELECT
15     cms.customerId,
16     c.customerName,
17     TO_CHAR(cms.saleMonth, 'YYYY-MM') AS saleMonthFormatted,
18     cms.monthlySales,
19     SUM(cms.monthlySales) OVER (PARTITION BY cms.customerId ORDER
20        BY cms.saleMonth ASC
21                                ROWS BETWEEN UNBOUNDED PRECEDING
22        AND CURRENT ROW) AS runningTotalSales
23 FROM
24     CustomerMonthlySales2022 cms
25 JOIN
26     Customers c ON cms.customerId = c.customerId
27 ORDER BY
28     cms.customerId, cms.saleMonth
29 LIMIT 200;
```

3. Indexes: CREATE INDEX idxSalesTransDateIncl ON SalesTransactions (transactionDate) INCLUDE (customerId, totalAmount);

3.4 Solution for Exercise OWA-4

```
1 WITH Sales2022 AS (  
2     SELECT  
3         st.productId,  
4         st.customerId,  
5         st.totalAmount,  
6         p.category,  
7         r.regionName  
8     FROM  
9         SalesTransactions st  
10    JOIN  
11        Products p ON st.productId = p.productId  
12    JOIN  
13        Customers c ON st.customerId = c.customerId  
14    JOIN  
15        Regions r ON c.regionId = r.regionId  
16    WHERE  
17        st.transactionDate >= '2022-01-01' AND st.transactionDate < '  
18        2023-01-01'  
19 ),  
20 CategoryRegionSales AS (  
21     SELECT  
22         category,  
23         regionName,  
24         SUM(totalAmount) AS categoryRegionTotalSales  
25     FROM  
26         Sales2022  
27     GROUP BY  
28         category, regionName  
29 ),  
30 RankedAndOverallSales AS (  
31     SELECT  
32         category,  
33         regionName,  
34         categoryRegionTotalSales,  
35         RANK() OVER (ORDER BY categoryRegionTotalSales DESC) AS  
36         overallRank,  
37         SUM(categoryRegionTotalSales) OVER (PARTITION BY regionName) AS  
38         totalRegionSales,  
39         SUM(categoryRegionTotalSales) OVER (PARTITION BY category) AS  
40         totalCategorySales  
41     FROM  
42         CategoryRegionSales  
43 )  
44 SELECT  
45     category,  
46     regionName,  
47     categoryRegionTotalSales,  
48     overallRank,  
49     ROUND((categoryRegionTotalSales * 100.0 / totalRegionSales)::  
50     NUMERIC, 2) AS pctOfRegionSales,  
51     ROUND((categoryRegionTotalSales * 100.0 / totalCategorySales)::  
52     NUMERIC, 2) AS pctOfCategorySales
```

```
47 FROM
48     RankedAndOverallSales
49 WHERE
50     categoryRegionTotalSales > 10000
51 ORDER BY
52     overallRank;
```